

# Adversarial Search

Larry Holder  
School of EECS  
Washington State University

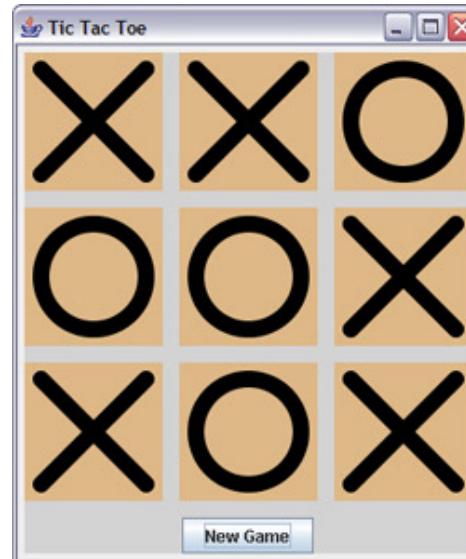
# Games

- ▶ Classic AI challenge
  - Easy to represent
  - Difficult to solve
- ▶ Perfect information (e.g., Chess, Checkers)
  - Fully observable and deterministic
- ▶ Imperfect information (e.g., Poker)
- ▶ Chance (e.g., Backgammon)



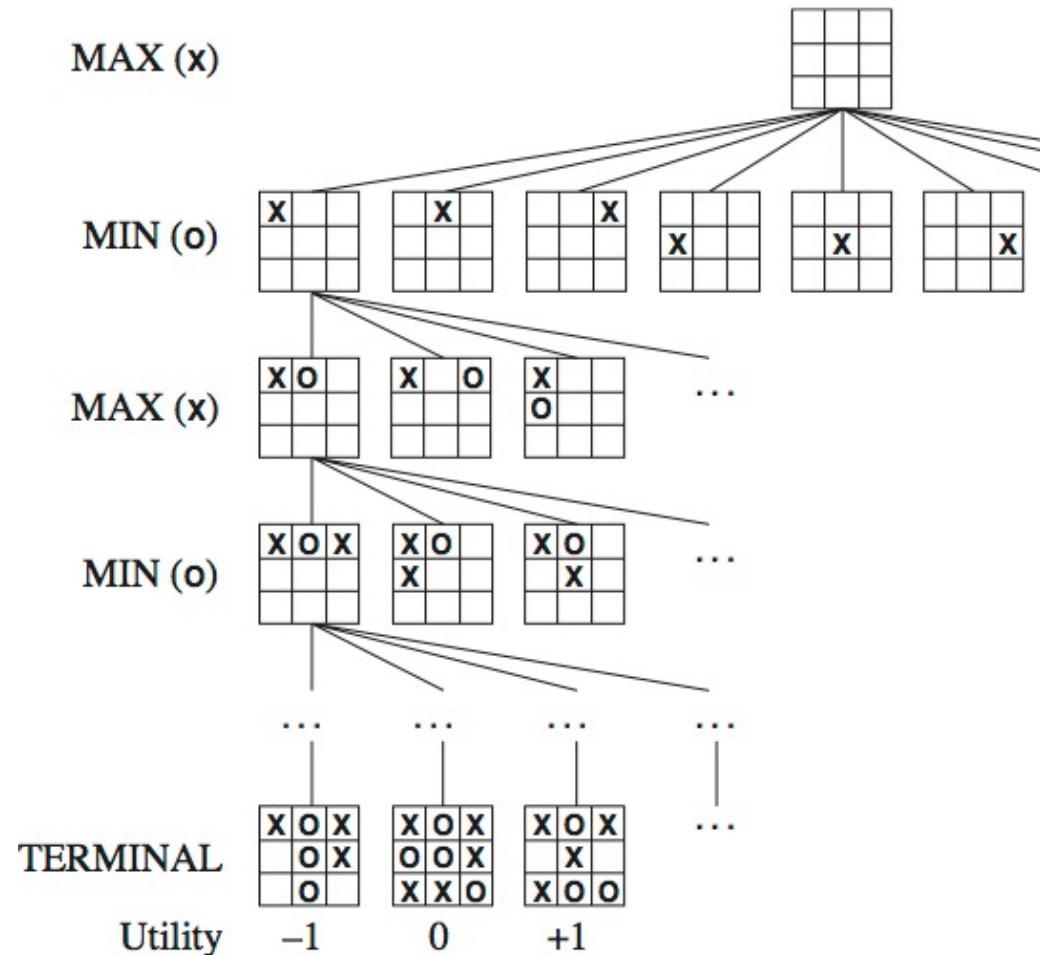
# Tic-Tac-Toe

- ▶ State space has about  $3^9 = 19,683$  nodes
- ▶ Average branching factor about 2
- ▶ Average game length about 8
- ▶ Search tree has about  $2^8 = 256$  nodes



# Game Tree

- ▶ MAX wants to maximize its outcome
- ▶ MIN wants to minimize its outcome
- ▶ Search tree refers to the search for a player's next move
- ▶ Terminal node
- ▶ Utility



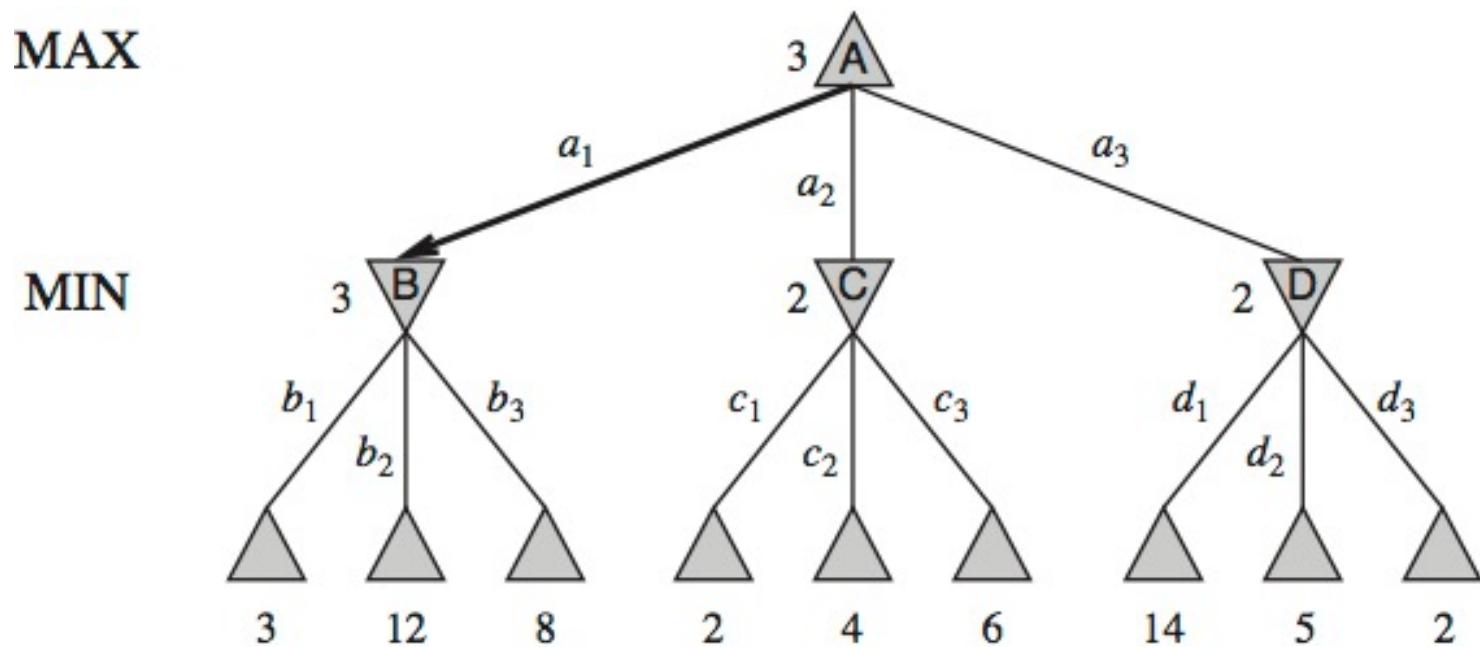
# Chess

- ▶ State space about  $10^{40}$  nodes
- ▶ Average branching factor about 35
- ▶ Average game length about 100 (50 moves per player)
- ▶ Search tree has about  $35^{100} = 10^{154}$  nodes



Garry Kasparov vs.  
IBM's Deep Blue  
(1997)

# Optimal Play



# Optimal Play

## ► Minimax value

- Best player can achieve assuming all players play optimally

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s) & \text{if TerminalTest}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MIN} \end{cases}$$

## ► Minimax decision

- Action that leads to minimax value

# Minimax Algorithm

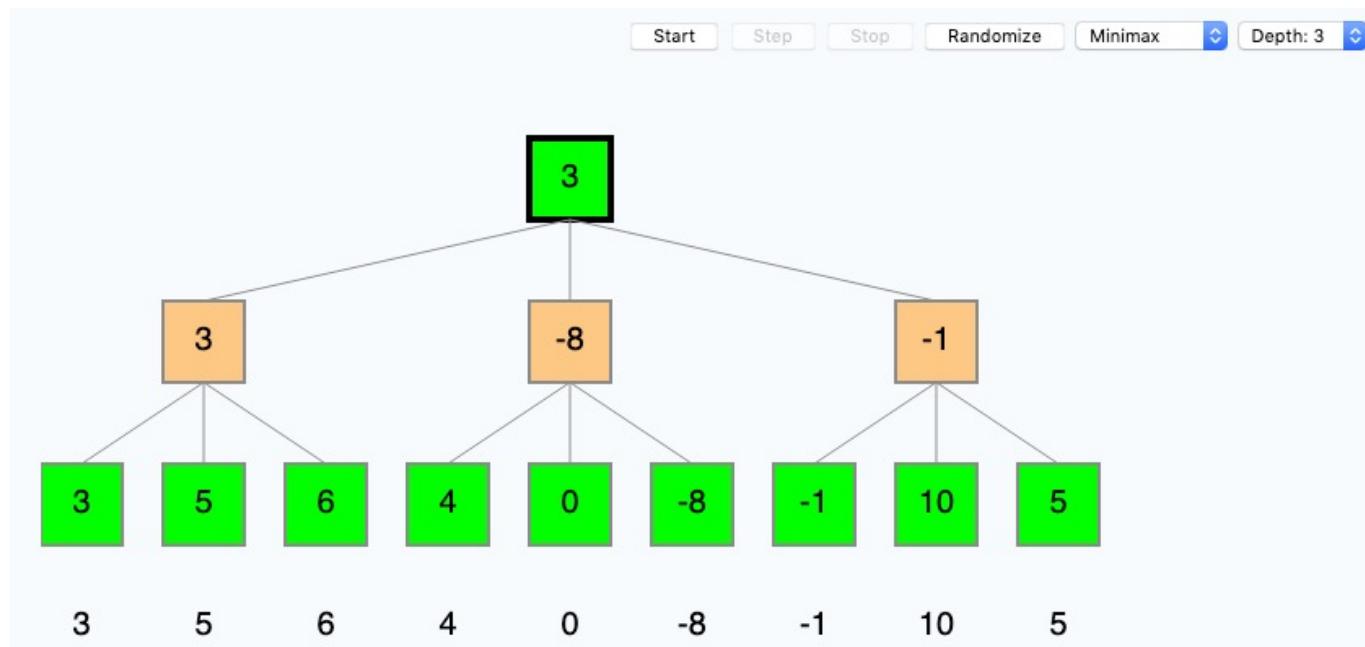
```
function MINIMAX-DECISION (state) returns an action  
  return arg maxa ∈ ACTIONS(state) MIN-VALUE(RESULT(state,a))
```

```
function MAX-VALUE (state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
  v ← -∞  
  for each a in ACTIONS(state) do  
    v ← MAX(v, MIN-VALUE(RESULT(state,a)))  
  return v
```

```
function MIN-VALUE (state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
  v ← ∞  
  for each a in ACTIONS(state) do  
    v ← MIN(v, MAX-VALUE(RESULT(state,a)))  
  return v
```

# Minimax Demo

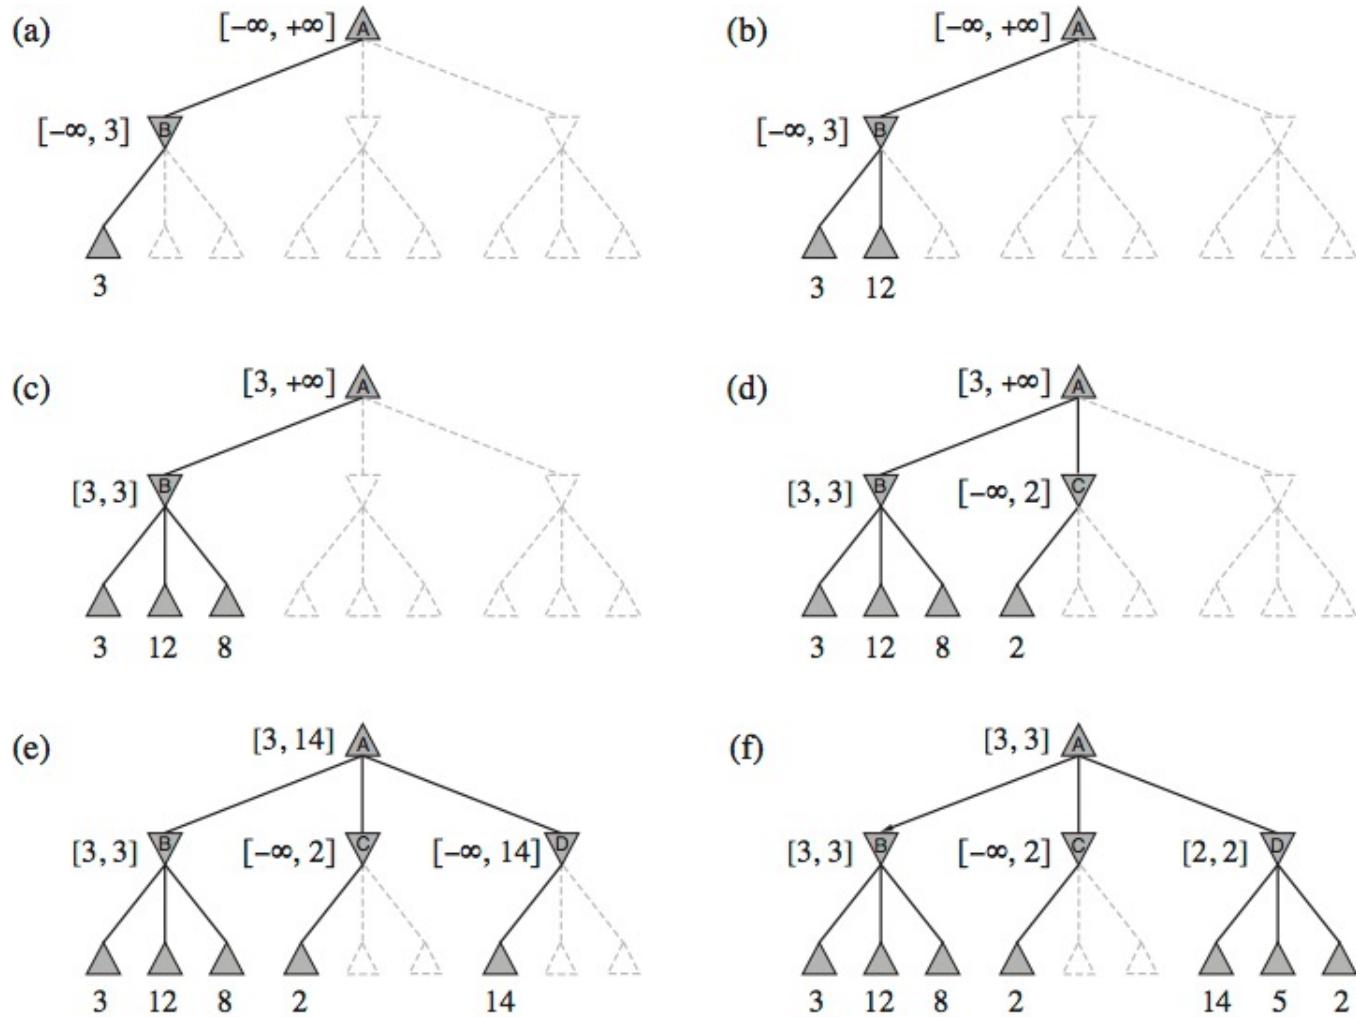
- ▶ [www.yosenspace.com/posts/computer-science-game-trees.html](http://www.yosenspace.com/posts/computer-science-game-trees.html)



# Minimax Algorithm

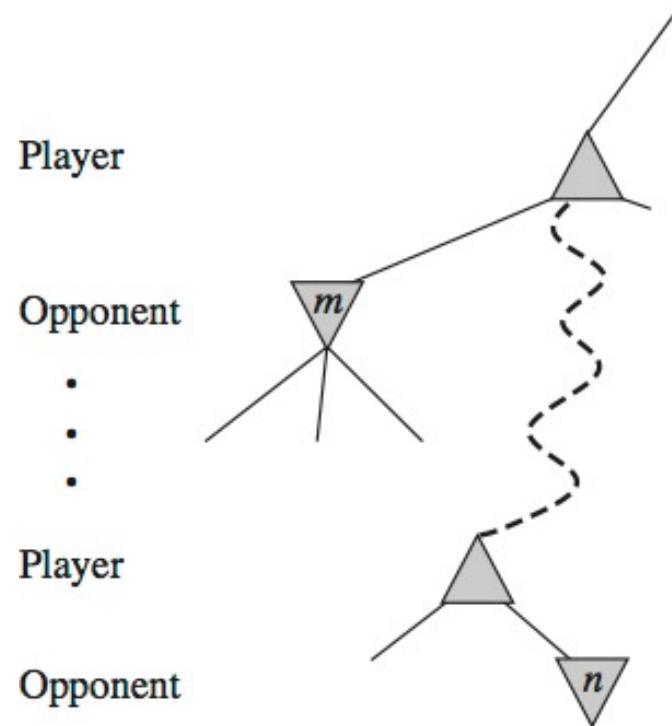
- ▶ Essentially depth-first search of game tree
- ▶ Time complexity:  $O(b^m)$ 
  - $m$  = maximum tree depth
  - $b$  = legal moves at each state
- ▶ Space complexity
  - Generates all actions:  $O(bm)$
  - Generates one action:  $O(m)$
- ▶ Practical?

# Pruning Search Tree



# Alpha–Beta Pruning

- ▶ Prune parts of the search tree that MAX and MIN would never choose
- ▶  $\alpha$  = value of best choice for MAX so far (highest value)
- ▶  $\beta$  = value of best choice for MIN so far (lowest value)
- ▶ Keep track of alpha  $\alpha$  and beta  $\beta$  during search



If  $m > n$ , Player will never move to  $n$ .

# Alpha-Beta Pruning

**function** ALPHA-BETA-SEARCH (*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\textit{state}, -\infty, +\infty)$

**return** the *action* in ACTIONS(*state*) with value  $v$

**function** MAX-VALUE (*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for each**  $a$  in ACTIONS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\textit{state}, a), \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

**return**  $v$

**function** MIN-VALUE (*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

**for each**  $a$  in ACTIONS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\textit{state}, a), \alpha, \beta))$

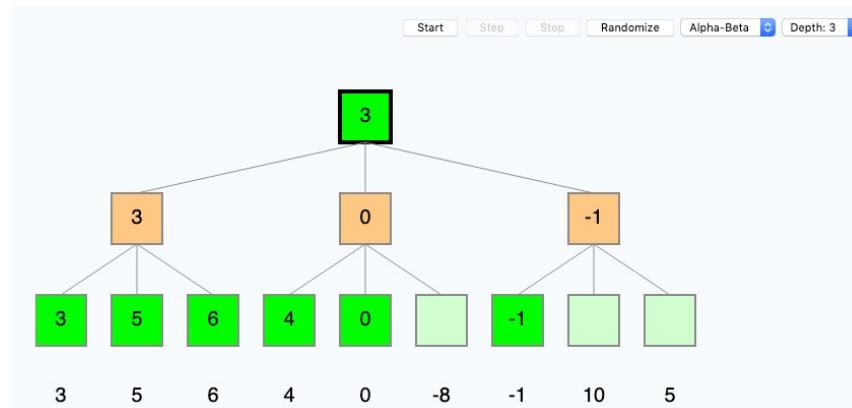
**if**  $v \leq \alpha$  **then return**  $v$

$\beta \leftarrow \text{MIN}(\beta, v)$

**return**  $v$

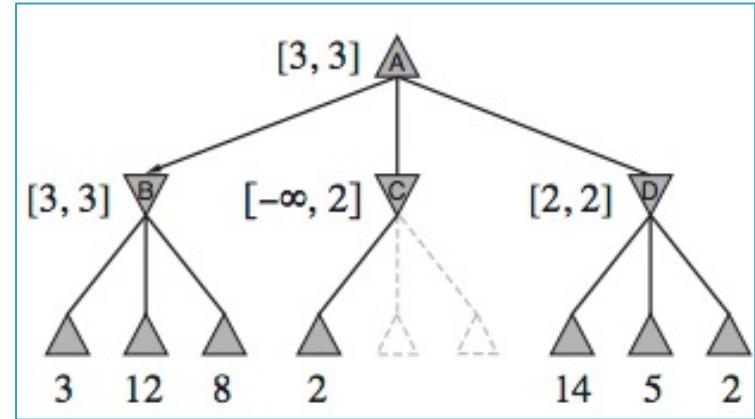
# Alpha-Beta Pruning Demo

- ▶ [www.yosenspace.com/posts/computer-science-game-trees.html](http://www.yosenspace.com/posts/computer-science-game-trees.html)



# Move Ordering

- ▶ ALPHA-BETA-SEARCH still  $O(b^m)$  worst case
- ▶ If order moves by value, then could prune maximally (always choose best move next)
  - Achieve  $O(b^{m/2})$  time
  - Branching factor  $b^{1/2}$
  - Chess: 35 → 6
  - But not practical
- ▶ Choosing moves randomly
  - Achieve  $O(b^{3m/4})$  average case
- ▶ Choosing moves based on impact
  - E.g., chess: captures, threats, forward, backward
  - Closer to  $O(b^{m/2})$



# Real-Time Game Play

- ▶ Minimax and Alpha-Beta search to terminal nodes
- ▶ Impractical for most games due to time limits
- ▶ Employ cutoff test to treat nodes as terminal nodes
- ▶ Heuristic evaluation function at these nodes to estimate utility
- ▶  $d$  = depth

$H\text{-Minimax}(s, d) =$

$$\begin{cases} \text{Eval}(s) & \text{if } \text{CutoffTest}(s, d) \\ \max_{a \in \text{Actions}(s)} H\text{-Minimax}(\text{Result}(s, a), d + 1) & \text{if } \text{Player}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} H\text{-Minimax}(\text{Result}(s, a), d + 1) & \text{if } \text{Player}(s) = \text{MIN} \end{cases}$$

# Real-Time Game Play

- ▶ Cutoff test
  - Depth-limit, iterative deepening until time's up
- ▶ Heuristic evaluation function EVAL(s)
  - Weighted combination of features

$$Eval(s) = \sum_{i=1}^n w_i f_i(s)$$

- E.g., chess
  - $f_1(s) = \# \text{pawns}$ ,  $w_1 = 1$
  - $f_4(s) = \# \text{bishops}$ ,  $w_4 = 3$
- Learn weights
- Learn features

# Go

- ▶ State space about  $10^{170}$  nodes
- ▶ Average branching factor about 250
- ▶ Average game length about 200 (100 moves per player)
- ▶ Search tree has about  $250^{200} = 10^{480}$  nodes



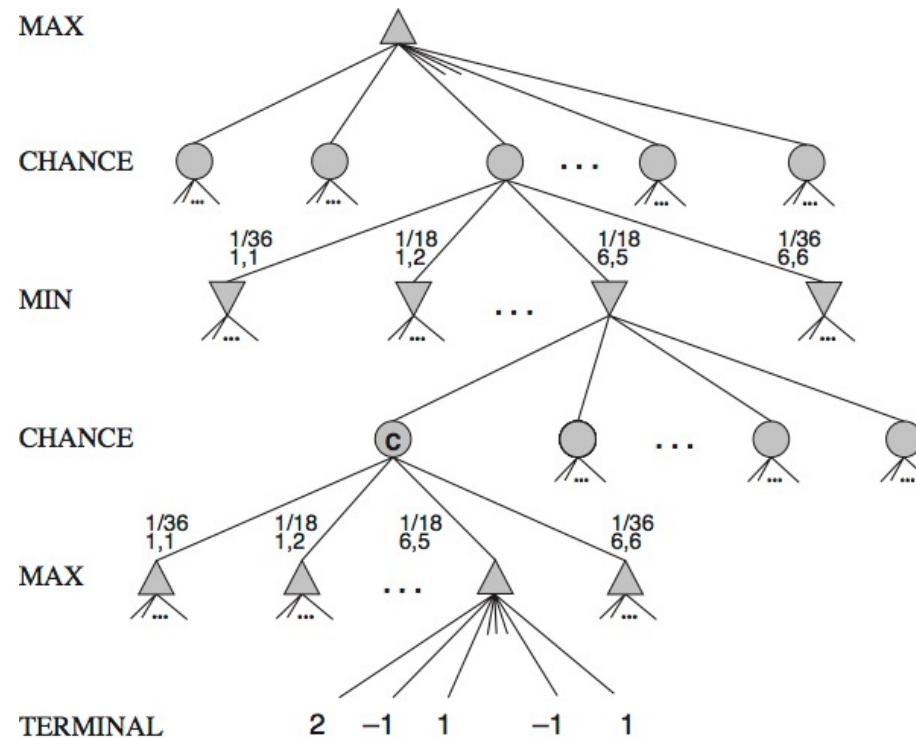
Lee Sodol vs.  
Google DeepMind's  
AlphaGo (2016)

[deepmind.com/research/alphago](http://deepmind.com/research/alphago)

# Stochastic Games



- ▶ Element of chance (e.g., dice roll)
- ▶ Include chance nodes in game tree
  - Branch to possible outcomes with their probabilities



# Stochastic Games

- ▶ Can't compute minimax values
- ▶ Can compute expected minimax values

$\text{ExpectiMinimax}(s) =$

$$\begin{cases} \text{Utility}(s) & \text{if } \text{TerminalTest}(s) \\ \max_{a \in \text{Actions}(s)} \text{ExpectiMinimax}(\text{Result}(s, a)) & \text{if } \text{Player}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{ExpectiMinimax}(\text{Result}(s, a)) & \text{if } \text{Player}(s) = \text{MIN} \\ \sum_r P(r) \text{ExpectiMinimax}(\text{Result}(s, r)) & \text{if } \text{Player}(s) = \text{CHANCE} \end{cases}$$

- $r$  represents possible chance event (e.g., dice roll)
- $\text{Result}(s, r) =$  state  $s$  with a particular outcome  $r$

# Stochastic Games

- ▶ Chance nodes increase branching factor
- ▶ Search time complexity  $O(b^m n^m)$ 
  - Where  $n$  is the number of chance outcomes
  - E.g., backgammon:  $n = 21$ ,  $b \approx 20$  (can be large)
  - Can only search a few moves ahead
- ▶ Estimate ExpectiMinimax values

# Partially Observable Games

- ▶ Can reason about all possible states of unknown information
- ▶ If  $P(s)$  represents probability of each unknown state  $s$ , then best move is:

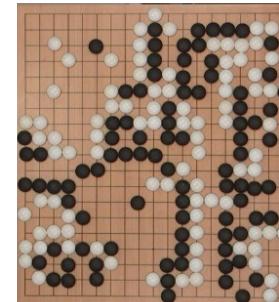
$$\arg \max_a \sum_s P(s) \text{Minimax}(\text{Result}(s, a))$$

- ▶ If  $|s|$  too large, take a random sample
  - Monte Carlo method



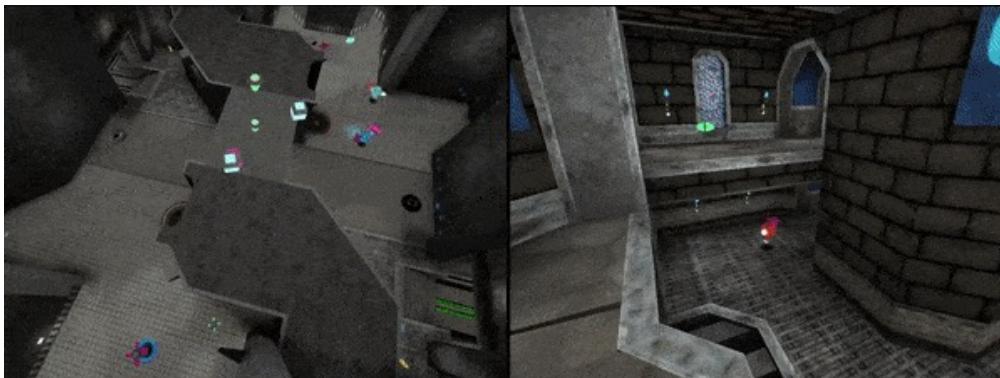
# State of the Art

- ▶ Checkers (solved, perfect play)
  - Chinook ([webdocs.cs.ualberta.ca/~chinook](http://webdocs.cs.ualberta.ca/~chinook))
  - Open/close database plus brute-force search
- ▶ Chess
  - Komodo ([komodochess.com](http://komodochess.com)) – proprietary
  - Stockfish ([stockfishchess.org](http://stockfishchess.org)) – open source
- ▶ Go
  - AlphaGo ([deepmind.com/research/alphago](http://deepmind.com/research/alphago))
  - Zen ([senseis.xmp.net/?ZenGoProgram](http://senseis.xmp.net/?ZenGoProgram))
- ▶ Backgammon
  - Extreme Gammon ([www.extremegammon.com](http://www.extremegammon.com))
  - GNU Backgammon ([www.gnu.org/software/gnubg](http://www.gnu.org/software/gnubg))
  - Neural network based evaluation function
- ▶ Poker
  - DeepStack ([www.deepstack.ai](http://www.deepstack.ai))
  - Pluribus ([ai.facebook.com/blog/pluribus-first-ai-to-beat-pros-in-6-player-poker](http://ai.facebook.com/blog/pluribus-first-ai-to-beat-pros-in-6-player-poker))



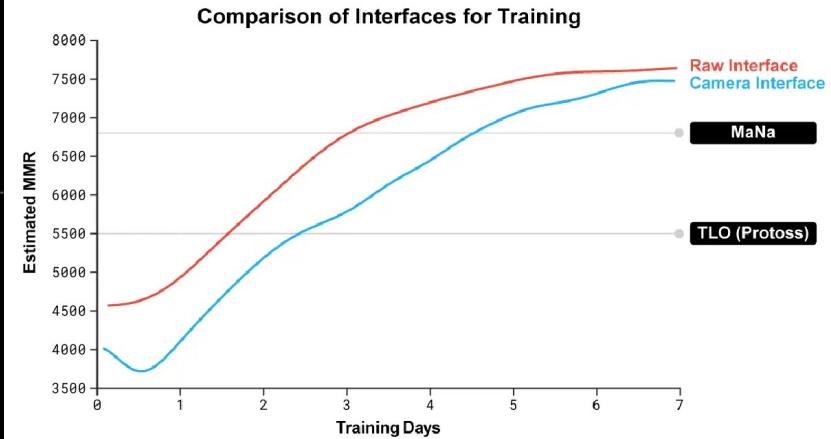
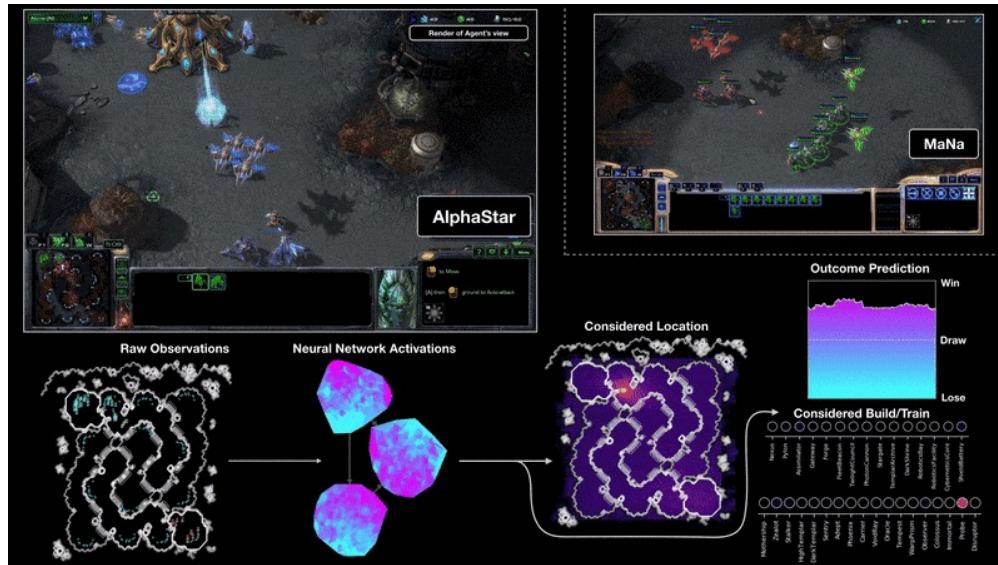
# State of the Art

- ▶ First-person shooter (FPS) games
  - DeepMind’s “For-The-Win” (FTW) Quake III agent
  - [deepmind.com/blog/article/capture-the-flag-science](https://deepmind.com/blog/article/capture-the-flag-science)



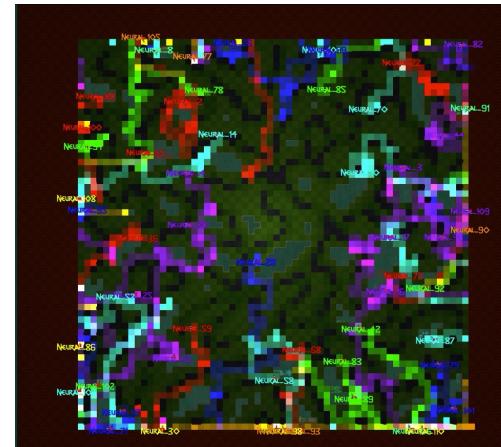
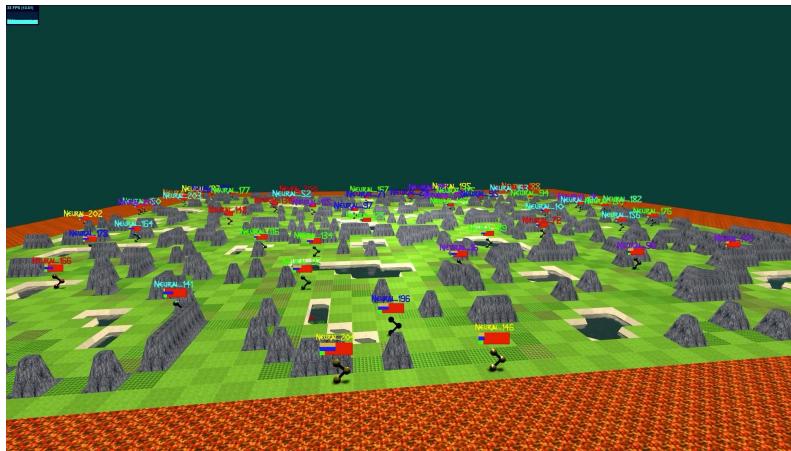
# State of the Art

- ▶ Real-Time Strategy (RTS) games
    - DeepMind’s AlphaStar masters StarCraft



# State of the Art

- ▶ Role-playing games (RPG/MMORPG)
- ▶ Neuro MMO
  - [openai.com/blog/neural-mmo](https://openai.com/blog/neural-mmo)



# Summary

- ▶ Adversarial search and games
- ▶ Minimax search
- ▶ Alpha–beta pruning
- ▶ Real–time issues
- ▶ Stochastic and partially observable games
- ▶ State of the art ...

Are there any games that humans can still beat computers?