

Uncertainty

Larry Holder
School of EECS
Washington State University

Uncertainty

- ▶ Sometimes the truth or falsity of facts in the world is unknown
- ▶ Sources of agent's uncertainty
 - Incompleteness of rules
 - Incorrectness of rules
 - Limited and ambiguous sensors
 - Imperfection and noise of actions
 - Unpredictable and dynamic nature of environment
 - Approximate nature of internal models and algorithms

Environment Properties

- ▶ Fully observable vs. **partially observable**
- ▶ Deterministic vs. **stochastic**
- ▶ Episodic vs. sequential
- ▶ Static vs. dynamic
- ▶ Discrete vs. continuous
- ▶ Single agent vs. multiagent

Wumpus World

- ▶ $P(\text{Pit}_{1,3}) = ?$
- ▶ $P(\text{Pit}_{2,2}) = ?$
- ▶ $P(\text{Pit}_{3,1}) = ?$

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2 B OK	2,2	3,2	4,2
1,1 OK	2,1 B OK	3,1	4,1

Rational Agent Approach

- ▶ Choose action A that maximizes expected utility
- ▶ I.e., maximizes $\text{Prob}(A) * \text{Utility}(A)$
- ▶ $\text{Prob}(A)$ = probability A will succeed
- ▶ $\text{Utility}(A)$ = value to agent of A's outcomes

$$\text{action} = \max_a P(a) * U(a)$$

Probability

- ▶ A probability model associates a numeric probability $P(w)$ with each possible world w , where

$$0 \leq P(w) \leq 1 \quad \text{and} \quad \sum_w P(w) = 1$$

Probability

- ▶ The probability $P(a)$ of a proposition ‘ a ’ is the sum of the probabilities of the worlds in which ‘ a ’ is true

$$P(a) = \sum_{w \in a} P(w)$$

Example

- ▶ What is the probability of the proposition that a two-dice roll totals 8?
- ▶ Consider the 36 possible outcomes of rolling two dice, each with probability $1/36$
 - (1, 1) (1, 2) (1, 3) (1, 4) (1, 5) (1, 6)
 - (2, 1) (2, 2) (2, 3) (2, 4) (2, 5) (2, 6)
 - (3, 1) (3, 2) (3, 3) (3, 4) (3, 5) (3, 6)
 - (4, 1) (4, 2) (4, 3) (4, 4) (4, 5) (4, 6)
 - (5, 1) (5, 2) (5, 3) (5, 4) (5, 5) (5, 6)
 - (6, 1) (6, 2) (6, 3) (6, 4) (6, 5) (6, 6)
- ▶ $P(\text{Total}=8) = P(\text{Die}_1=2 \wedge \text{Die}_2=6) + P(\text{Die}_1=3 \wedge \text{Die}_2=5) + \dots + P(\text{Die}_1=6 \wedge \text{Die}_2=2) = ?$

Probability

- ▶ An unconditional or prior probability is the degree of belief without any other information
 - E.g., $P(\text{Total}=8)$
- ▶ A conditional or posterior probability is the degree of belief given some evidence
 - E.g., given that the first die is a 2, what is the probability that the total will be 8?
 - $P(\text{Total}=8 \mid \text{Die}_1=2) = ?$
 - $P(\text{Total}=8 \mid \text{Die}_1=1) = ?$

Conditional Probability

$$P(a | b) = \frac{P(a \wedge b)}{P(b)} \quad \text{assumes} \quad P(b) > 0$$

- ▶ Fraction of worlds in which ‘a’ and ‘b’ are true out of the worlds in which ‘b’ is true
- ▶ $P(\text{Total}=8 | \text{Die}_1=2)$
 $= P(\text{Total}=8 \wedge \text{Die}_1=2) / P(\text{Die}_1=2) = ?$
- ▶ Product rule: $P(a \wedge b) = P(a | b) P(b)$

Probability Theory

- ▶ Random variable: a variable in probability theory
 - E.g., Total, Die₁, Die₂
- ▶ Domain: set of possible values for a random variable
 - E.g., domain(Die₁) = {1,2,3,4,5,6}
- ▶ Probability distribution: probability of each value in a random variable's domain
 - E.g., $P(\text{Die}_1=1)=1/6$, ..., $P(\text{Die}_1=6)=1/6$
 - Or, $\mathbf{P}(\text{Die}_1)=\langle 1/6, 1/6, 1/6, 1/6, 1/6, 1/6 \rangle$

Note: Boldfaced

Probability Theory


- ▶ Conditional probability distribution
 - $P(X | Y)$ gives the values of $P(X=x_i | Y=y_j)$ for all possible i,j pairs
- ▶ Joint probability distribution
 - $P(X,Y)$ denotes the probabilities of all possible combinations of X and Y
 - E.g., $P(\text{Die}_1, \text{Die}_2)$ gives values for $P(\text{Die}_1=1, \text{Die}_2=1)$, $P(\text{Die}_1=1, \text{Die}_2=2)$, ..., $P(\text{Die}_1=6, \text{Die}_2=6)$
 - E.g., $P(\text{Die}_1=5, \text{Die}_2=1)$ abbreviated $P(5,1)$
- ▶ Full joint probability distribution
 - Joint probability distribution over all random variables in the world

Wumpus World

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2 B OK	2,2	3,2	4,2
1,1 OK	2,1 B OK	3,1	4,1


More Probability Axioms

- ▶ $P(\neg a) = 1 - P(a)$
- ▶ $P(a \vee b) = P(a) + P(b) - P(a \wedge b)$



Probabilistic Inference

- ▶ Approach #1: Just ask Spock...



“Errand of Mercy” (1967)

Probabilistic Inference

- Given full joint probability distribution, we can answer any question about probabilities
- Example: Tooth World
 - Random variables: Toothache, Cavity, Catch
 - Domain: true, false

	Toothache = true		Toothache = false	
	Catch = true	Catch = false	Catch = true	Catch = false
Cavity = true	.108	.012	.072	.008
Cavity = false	.016	.064	.144	.576

Probabilistic Inference

- ▶ Answer questions (perform probabilistic inference) by summing probabilities
- ▶ Examples
 - $P(\text{toothache}) = ?$
 - $P(\text{cavity} \wedge \text{toothache}) = ?$
 - $P(\text{cavity} \mid \text{toothache}) = ?$

	toothache		$\neg\text{toothache}$	
	catch	$\neg\text{catch}$	catch	$\neg\text{catch}$
cavity	.108	.012	.072	.008
$\neg\text{cavity}$.016	.064	.144	.576

Note: For Boolean random variables, write $X=\text{true}$ as x , and $X=\text{false}$ as $\neg x$.

Probabilistic Inference

- ▶ Marginalization is the process of finding the probability distribution over a subset of variables Y :

$$P(Y) = \sum_{z \in Z} P(Y, z)$$

- For example:

$$P(Cavity) = \sum_{z \in \{Catch, Toothache\}} P(Cavity, z)$$

z in {(catch, tooth), (catch, \neg tooth),
 $(\neg$ catch, tooth), (\neg catch, \neg tooth)}

- ▶ Conditioning:

$$P(Y) = \sum_{z \in Z} P(Y | z)P(z)$$

Probabilistic Inference

▶ Conditional probabilities

$$\begin{aligned} P(cavity | toothache) &= \frac{P(cavity \wedge toothache)}{P(toothache)} \\ &= \frac{0.108 + 0.012}{0.108 + 0.012 + 0.016 + 0.064} = 0.6 \end{aligned}$$

$$\begin{aligned} P(\neg cavity | toothache) &= \frac{P(\neg cavity \wedge toothache)}{P(toothache)} \\ &= \frac{0.016 + 0.064}{0.108 + 0.012 + 0.016 + 0.064} = 0.4 \end{aligned}$$

▶ Note that denominator is the same in both (normalization constant)

Probabilistic Inference

- ▶ Don't need normalization constant to compute probability distribution
 - For example
 - $P(\text{cavity} \mid \text{toothache}) = \alpha P(\text{cavity} \wedge \text{toothache}) = \alpha 0.12$
 - $P(\neg\text{cavity} \mid \text{toothache}) = \alpha P(\neg\text{cavity} \wedge \text{toothache}) = \alpha 0.08$
 - Since these have to sum to 1, $\alpha = 1/0.2 = 5$
 - Also means $P(\text{toothache}) = 1/5 = 0.2$, but we didn't need to know this
 - Also means we can determine which is more likely (cavity or \neg cavity) without knowing $P(\text{toothache})$

Probabilistic Inference

- ▶ General rule
- ▶ Want to know the probability distribution over X given observed variables e (evidence) and unobserved variables y

$$P(X | e) = \alpha P(X, e) = \alpha \sum_y P(X, e, y)$$

Wumpus World

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2 B OK	2,2	3,2	4,2
1,1 OK	2,1 B OK	3,1	4,1

Independence

- ▶ Full joint probability distributions are typically infeasible to write down
 - E.g., 2^n table entries for n Boolean variables
- ▶ If we know some variables are independent of others, then we can decompose the full table into smaller tables
- ▶ If two variables X and Y are independent, then
 - $P(X, Y) = P(X)P(Y)$
 - $P(X|Y) = P(X)$
 - $P(Y|X) = P(Y)$

Independence

- ▶ For example, add Weather variable to Tooth World
 - $\text{domain}(\text{Weather}) = \{\text{sunny}, \text{cloudy}, \text{rain}, \text{snow}\}$
 - 8 entry table now a 32 entry table
- ▶ Assuming your teeth don't affect the weather
 - I.e., Weather is independent of Toothache, Catch and Cavity
 - So, $P(\text{Weather}, \text{Toothache}, \text{Catch}, \text{Cavity}) = P(\text{Weather}) * P(\text{Toothache}, \text{Catch}, \text{Cavity})$
 - E.g., $P(\text{Weather}=\text{cloudy}, \text{toothache}, \text{catch}, \text{cavity}) = P(\text{Weather}=\text{cloudy}) * P(\text{toothache}, \text{catch}, \text{cavity})$
- ▶ Full joint distribution described by two tables of 8 and 4 entries

Bayes Rule

Thomas Bayes
(1701–1761)



$$P(b | a) = \frac{P(a | b)P(b)}{P(a)}$$

$$\mathbf{P}(Y | X) = \frac{\mathbf{P}(X | Y)\mathbf{P}(Y)}{\mathbf{P}(X)}$$

- ▶ Foundational rule of probabilistic reasoning
- ▶ Turns a diagnostic question into a causal one

$$P(\text{cause} | \text{effect}) = \frac{P(\text{effect} | \text{cause})P(\text{cause})}{P(\text{effect})}$$

- ▶ In general, given some evidence e :

$$\mathbf{P}(Y | X, e) = \frac{\mathbf{P}(X | Y, e)\mathbf{P}(Y | e)}{\mathbf{P}(X | e)}$$

Bayes Rule

- ▶ Example: Diagnosing cancer
- ▶ Given
 - 1% of population has cancer
 - Test has 20% false positive rate
 - Test has 10% false negative rate
 - Patient tests positive for cancer
 - What is the probability that patient has cancer?

Bayes Rule

- ▶ Example: Diagnosing cancer
- ▶ Variables
 - Diagnosis $\in \{\text{cancer}, \text{healthy}\}$
 - Test $\in \{\text{pos}, \text{neg}\}$
- ▶ Given
 - $P(\text{cancer})=0.01, P(\text{healthy})=0.99$
 - $P(\text{neg}|\text{cancer})=0.1, P(\text{pos}|\text{cancer})=0.9$
 - $P(\text{pos}|\text{healthy})=0.2, P(\text{neg}|\text{healthy})=0.8$
- ▶ $P(\text{cancer}|\text{pos})=?$
- ▶ $P(\text{cancer}|\text{pos}) = P(\text{cancer} \wedge \text{pos}) / P(\text{pos}) = ?$

Bayes Rule

► Applying Bayes rule...

$$P(cancer | pos) = \frac{P(pos | cancer)P(cancer)}{P(pos)}$$

- We know $P(pos | cancer) = 0.9$ and $P(cancer) = 0.01$
- Can compute $P(pos)$ by marginalization
- $P(pos) = ?$

- $P(cancer | pos) = ?$

Bayes Rule

- ▶ Could also compute $P(\text{pos})$ via normalization
- ▶ $P(\text{cancer}|\text{pos}) = \alpha P(\text{pos}|\text{cancer})P(\text{cancer}) =$
- ▶ $P(\text{healthy}|\text{pos}) = \alpha P(\text{pos}|\text{healthy})P(\text{healthy}) =$
- ▶ $\alpha =$
- ▶ $P(\text{pos}) = 1 / \alpha =$
- ▶ Many times, $P(\text{effect}|\text{cause})$ is easier to determine than $P(\text{cause}|\text{effect})$

Monty Hall Problem

- ▶ C_i = car behind Door i
- ▶ O_i = opens Door i
- ▶ Pick Door 2
- ▶ Opens Door 3



"Numb3rs" S1E13, 2005.

- ▶ Should you change your guess?
- ▶ $P(C_1 \mid O_3) = ?$

Combining Evidence in Bayes Rule

- ▶ How to compute $P(a | b \wedge c \wedge \dots)$?
 - E.g., $P(\text{Cavity} | \text{toothache} \wedge \text{catch}) = ?$
- ▶ Easy if we have full joint probability distribution
 - E.g., $P(\text{Cavity} | \text{toothache} \wedge \text{catch}) = \alpha \langle P(\text{cavity} \wedge \text{toothache} \wedge \text{catch}), P(\neg \text{cavity} \wedge \text{toothache} \wedge \text{catch}) \rangle = ?$
- ▶ Or, using Bayes rule: $\alpha [P(b \wedge c \wedge \dots | a) P(a)]$
 - E.g., $\alpha P(\text{toothache} \wedge \text{catch} | \text{Cavity}) P(\text{Cavity})$
 - Still need to know many probabilities

Combining Evidence in Bayes Rule

- ▶ If b, c, \dots are “caused” by ‘ a ’, but not by each other, then
 - $P(b \wedge c \wedge \dots | a) = P(b | a) P(c | a) \dots$
 - $P(a | b \wedge c \wedge \dots) = \alpha P(b \wedge c \wedge \dots | a) P(a)$
 - $P(a | b \wedge c \wedge \dots) = \alpha P(b | a) P(c | a) \dots P(a)$
- ▶ I.e., b, c, \dots are independent given ‘ a ’
- ▶ Two propositions (effects) are conditionally independent if they are independent given a third proposition (cause)
- ▶ Now we only need to know the probabilities of individual effects given the cause $P(b | a)$ and the prior probability of the cause $P(a)$

Combining Evidence in Bayes Rule

- ▶ Example (assuming toothache and catch are conditionally independent given Cavity)
 - $P(\text{toothache} \wedge \text{catch} | \text{Cavity}) = P(\text{toothache} | \text{Cavity}) P(\text{catch} | \text{Cavity})$
- ▶ Using this in Bayes rule
 - $P(\text{Cavity} | \text{toothache} \wedge \text{catch}) = \alpha P(\text{toothache} \wedge \text{catch} | \text{Cavity}) P(\text{Cavity})$
 - $P(\text{Cavity} | \text{toothache} \wedge \text{catch}) = \alpha P(\text{toothache} | \text{Cavity}) P(\text{catch} | \text{Cavity}) P(\text{Cavity})$

Combining Evidence in Bayes Rule

- ▶ In general, if random variables X and Y are conditionally independent given Z
 - $P(X, Y | Z) = P(X | Z) P(Y | Z)$
 - $P(X | Y, Z) = P(X | Z)$
 - $P(Y | X, Z) = P(Y | Z)$

Naïve Bayes

- ▶ Want $P(\text{Cause} \mid \text{Effect}_1, \text{Effect}_2, \dots, \text{Effect}_n)$
- ▶ Assume $\text{Effect}_1, \text{Effect}_2, \dots, \text{Effect}_n$ are conditionally independent given Cause
- ▶ Applying Bayes Rule
 - $P(\text{Cause} \mid \text{Effect}_1, \text{Effect}_2, \dots, \text{Effect}_n) = \alpha P(\text{Effect}_1, \text{Effect}_2, \dots, \text{Effect}_n \mid \text{Cause}) P(\text{Cause}) = \alpha P(\text{Cause}) \prod_i P(\text{Effect}_i \mid \text{Cause})$

Naïve Bayes

- ▶ Be careful
 - Conditional independence rarely absolutely true
 - E.g., Cold weather makes my teeth hurt
 - Probabilities estimated from data


Wumpus World Revisited

- ▶ $P(\text{Pit}_{1,3}) = ?$
- ▶ $P(\text{Pit}_{2,2}) = ?$
- ▶ $P(\text{Pit}_{3,1}) = ?$


1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2 B OK	2,2	3,2	4,2
1,1 OK	2,1 B OK	3,1	4,1

Wumpus World Revisited


- ▶ “query” = $\text{Pit}_{1,3}$
- ▶ “frontier” = { $\text{Pit}_{2,2}$, $\text{Pit}_{3,1}$ }
- ▶ “other” = other 10 pit variables
- ▶ “known” = $\neg \text{pit}_{1,1} \wedge \neg \text{pit}_{1,2} \wedge \neg \text{pit}_{2,1}$
- ▶ “breeze” = $\neg \text{breeze}_{1,1} \wedge \text{breeze}_{1,2} \wedge \text{breeze}_{2,1}$
- ▶ Note: “breeze” is conditionally independent of “other” given “known”, “frontier” and “query”



Wumpus World Revisited



Wumpus World Revisited



Wumpus World Revisited

- And given independence of pits $P_{i,j}$
- $P(P_{1,3} | \text{known, breeze}) = \alpha P(P_{1,3}) \sum_{\text{frontier}} P(\text{breeze} | \text{known}, P_{1,3}, \text{frontier}) P(\text{frontier}) = \alpha \langle 0.2(0.04+0.16+0.16), 0.8(0.04+0.16) \rangle = \langle 0.31, 0.69 \rangle$

1,3	
1,2 B OK	2,2 OK
1,1 OK	2,1 B OK

$$0.2 \times 0.2 = 0.04$$

1,3	
1,2 B OK	2,2 OK
1,1 OK	2,1 B OK

$$0.2 \times 0.8 = 0.16$$

$P_{1,3} = \text{true}$

1,3	
1,2 B OK	2,2 OK
1,1 OK	2,1 B OK

$$0.8 \times 0.2 = 0.16$$

1,3	
1,2 B OK	2,2 OK
1,1 OK	2,1 B OK

$$0.2 \times 0.2 = 0.04$$

$P_{1,3} = \text{false}$

1,3	
1,2 B OK	2,2 OK
1,1 OK	2,1 B OK

$$0.2 \times 0.8 = 0.16$$

Wumpus World Revisited

- ▶ So, $P(P_{1,3}=\text{true}) = P(P_{3,1}=\text{true}) = 0.31$
- ▶ $P(P_{2,2}=\text{true}) = 0.86$
- ▶ Probabilistic agent “knows more” than logical agent

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2 B OK	2,2	3,2	4,2
1,1 OK	2,1 B OK	3,1	4,1

Summary: Uncertainty

- ▶ Probability theory allows us to reason about uncertainty
- ▶ Bayes rule allows us to change diagnostic questions to causal questions
- ▶ Conditional independence allows us to simplify complexity
- ▶ Probabilistic agent can outperform logical agent in partially observable and stochastic environments