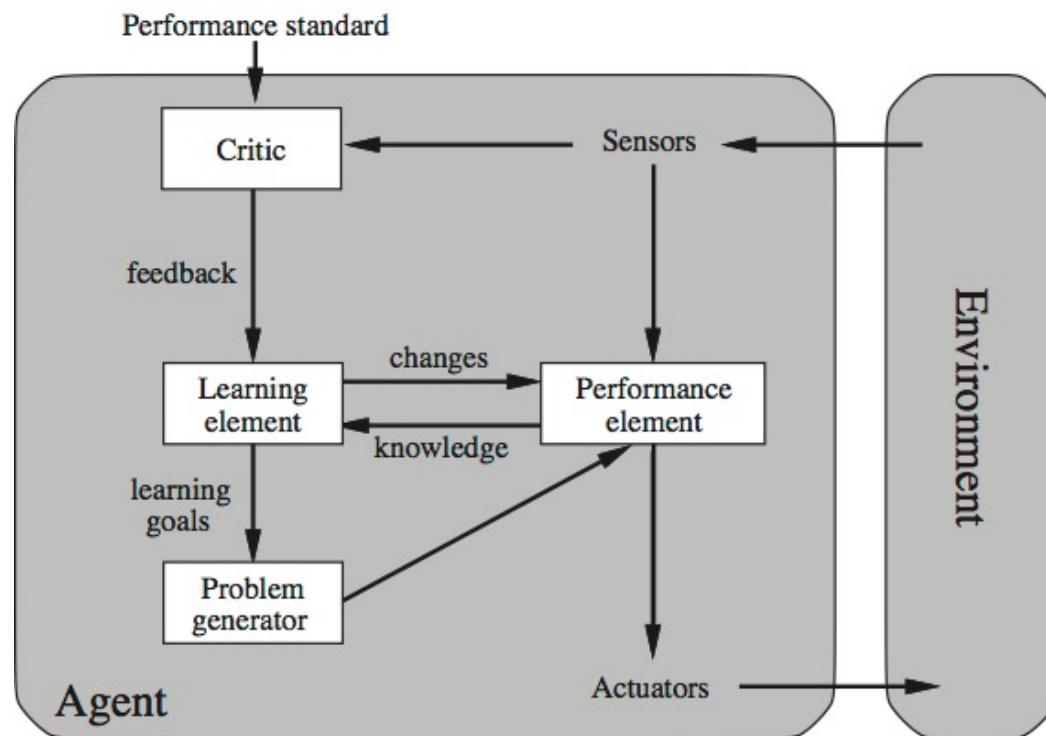


Learning

Larry Holder
School of EECS
Washington State University

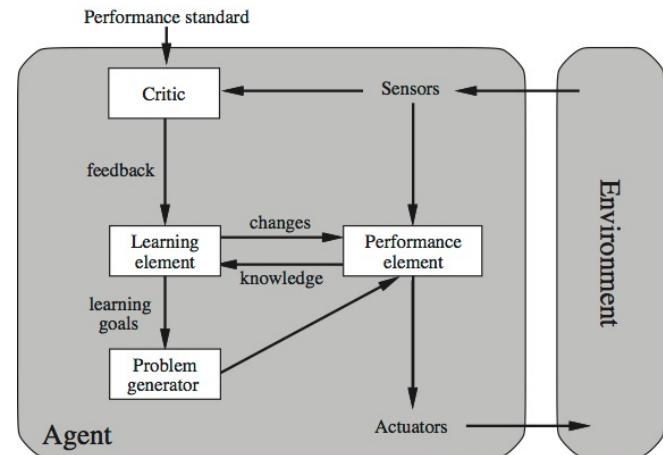
Learning Agent

- An agent that improves its performance at some task through experience



Details, details

- ▶ How is knowledge represented?
- ▶ How is performance measured and critiqued?
- ▶ How is experience represented and obtained?

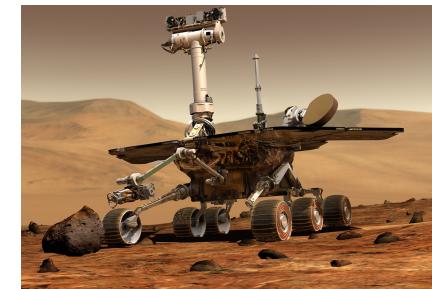
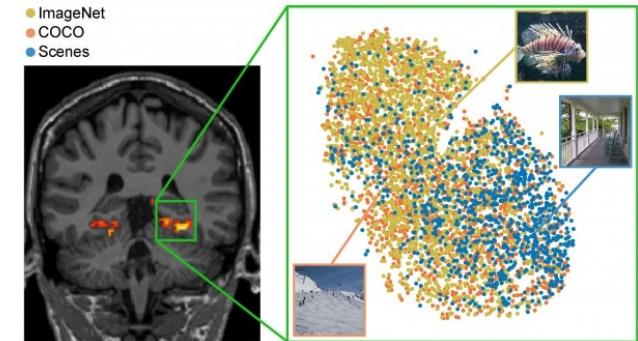
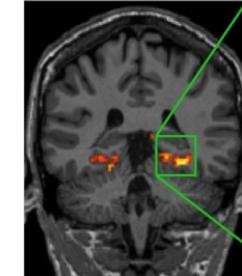


Why Do Machine Learning?

- ▶ Automated knowledge acquisition
- ▶ Discover new knowledge
- ▶ Understand human learning
- ▶ Agents need to adapt to unknown, dynamic environments



● ImageNet
● COCO
● Scenes



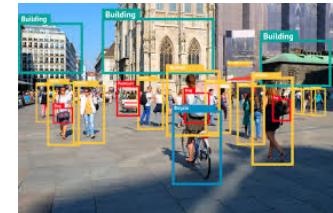
"If you invent a breakthrough in artificial intelligence, so machines can learn, that is worth ten Microsofts."

Bill Gates, 2004

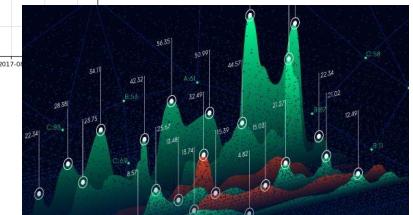
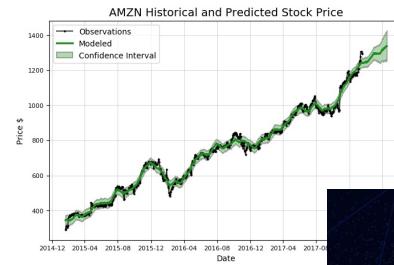


Applications

- ▶ Medical diagnosis
- ▶ Autonomous control (planes, trains, automobiles, robots, ...)
- ▶ Perception (speech, language, images, video)
- ▶ Recommendations (Amazon, Netflix)
- ▶ Prediction (business, financial, environment, health, energy, security, ...)
- ▶ Fraud/intrusion detection
- ▶ ...

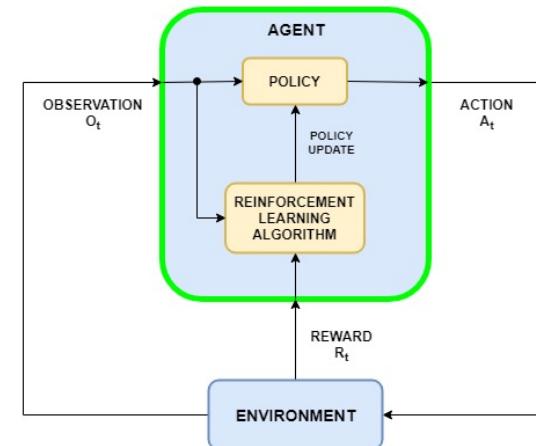
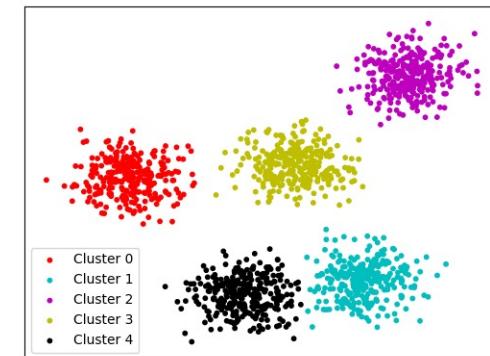
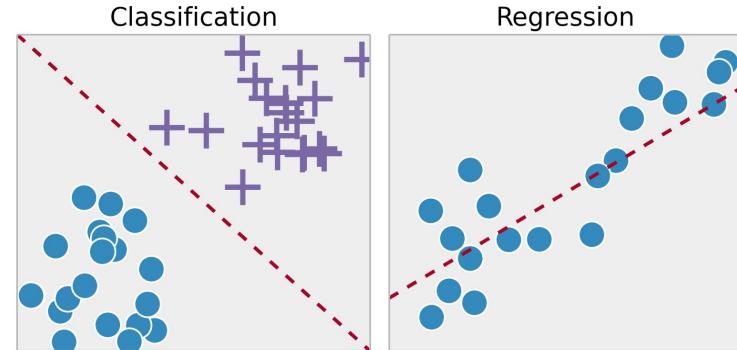


Customers who viewed this item also viewed



Approaches

- ▶ Supervised Learning
 - Classification
 - Regression
- ▶ Unsupervised Learning
 - Clustering
- ▶ Reinforcement Learning

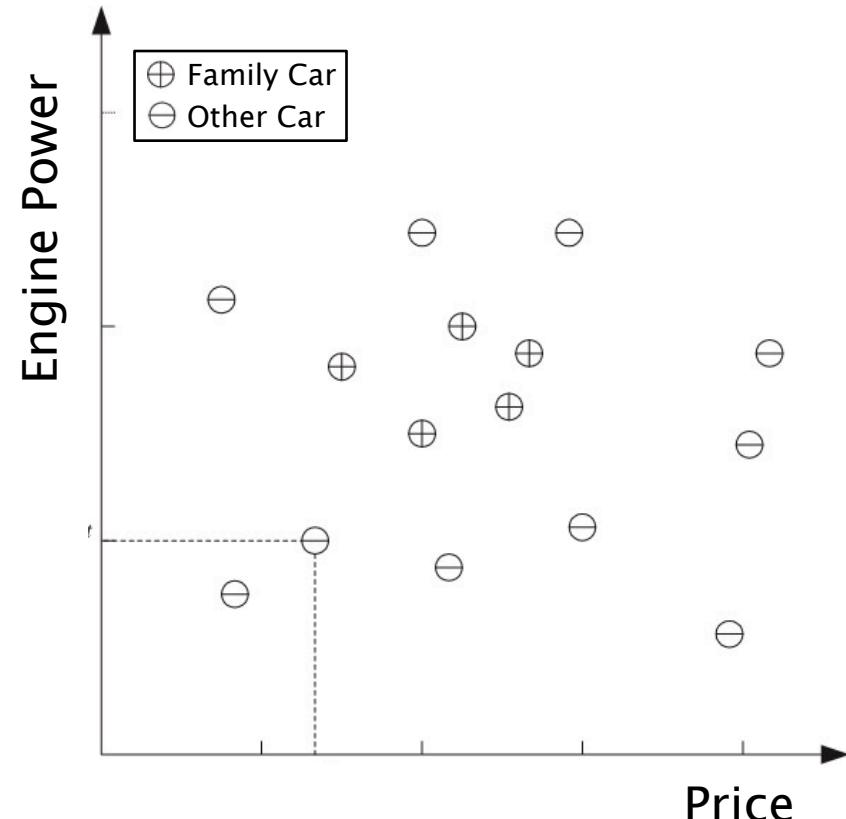


Supervised Learning

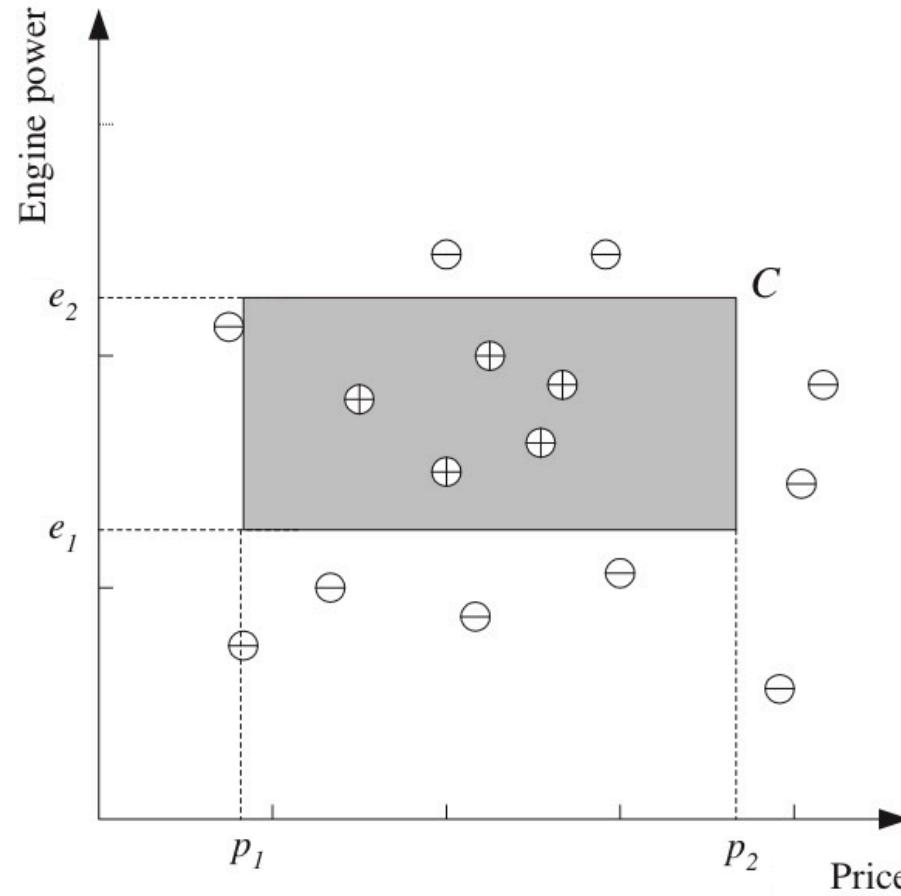
- ▶ Given training set of N examples of $y=f(x)$
 - $(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)$
- ▶ Find hypothesis h that approximates f

Example: “Family Car”

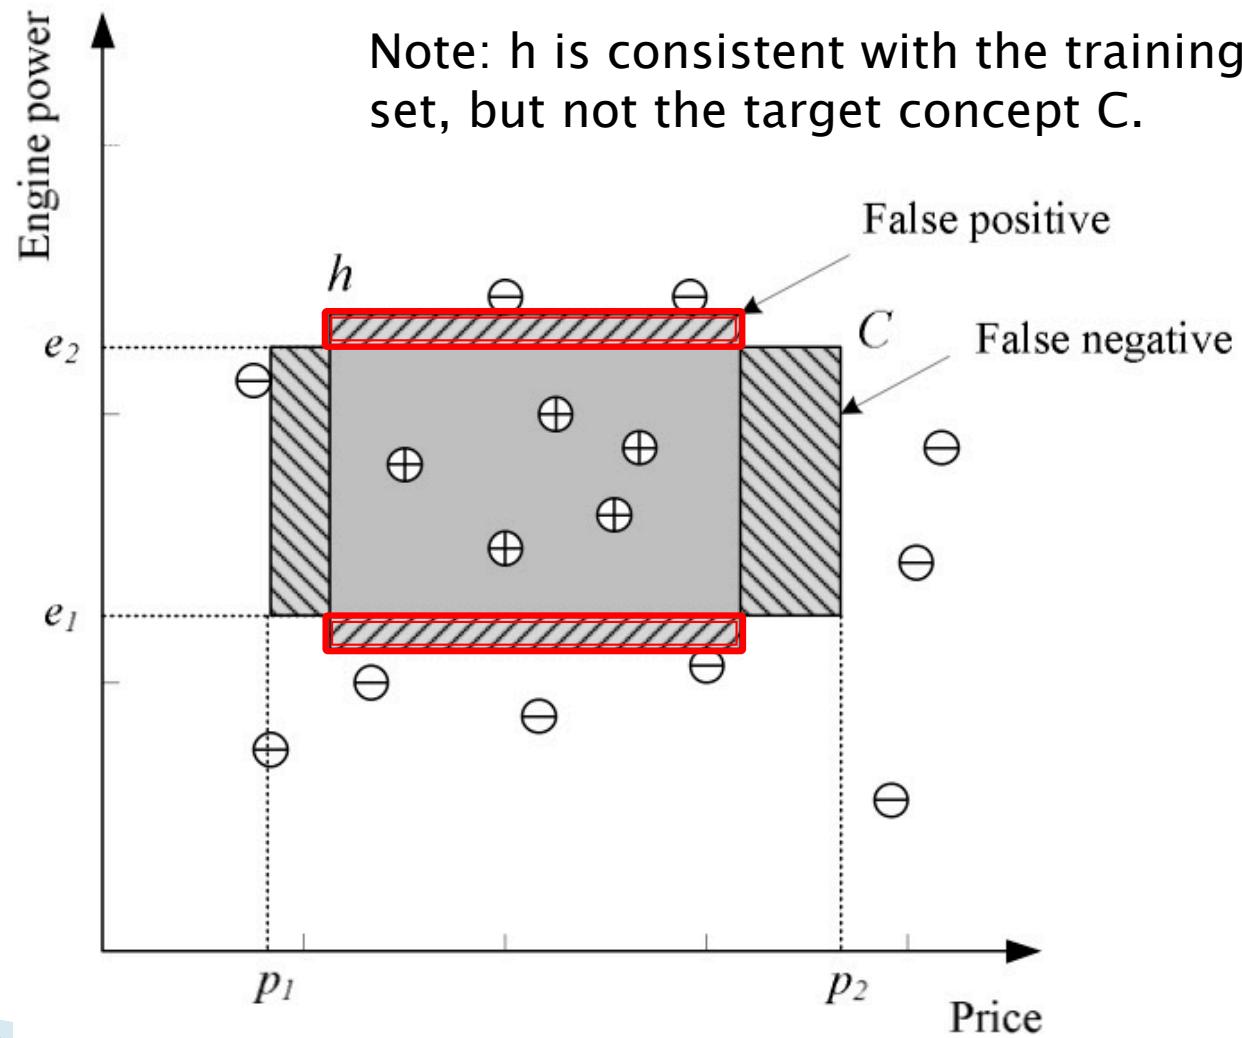
- ▶ Learn to classify cars into one of two classes: “family car” or “other”
- ▶ Each car is represented by two features: “engine power” and “price”
- ▶ Given several training examples of already-classified cars
- ▶ Output classifier that accurately classifies all cars



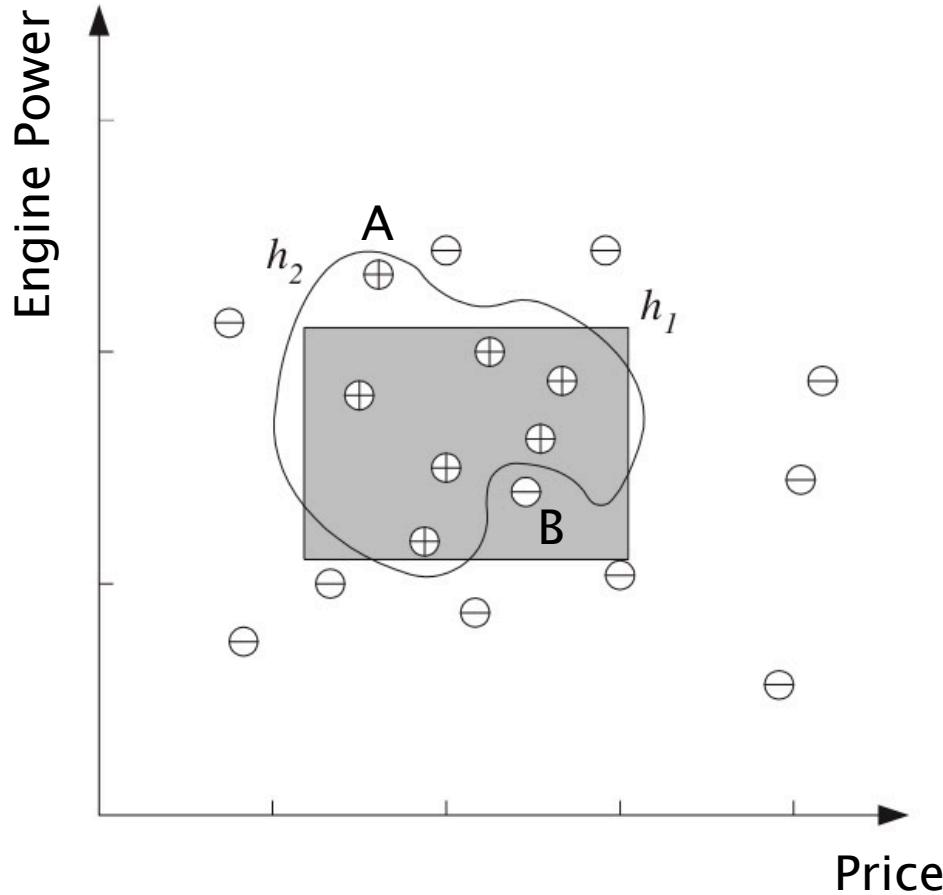
“Family Car” Target Concept



Hypothesis Error



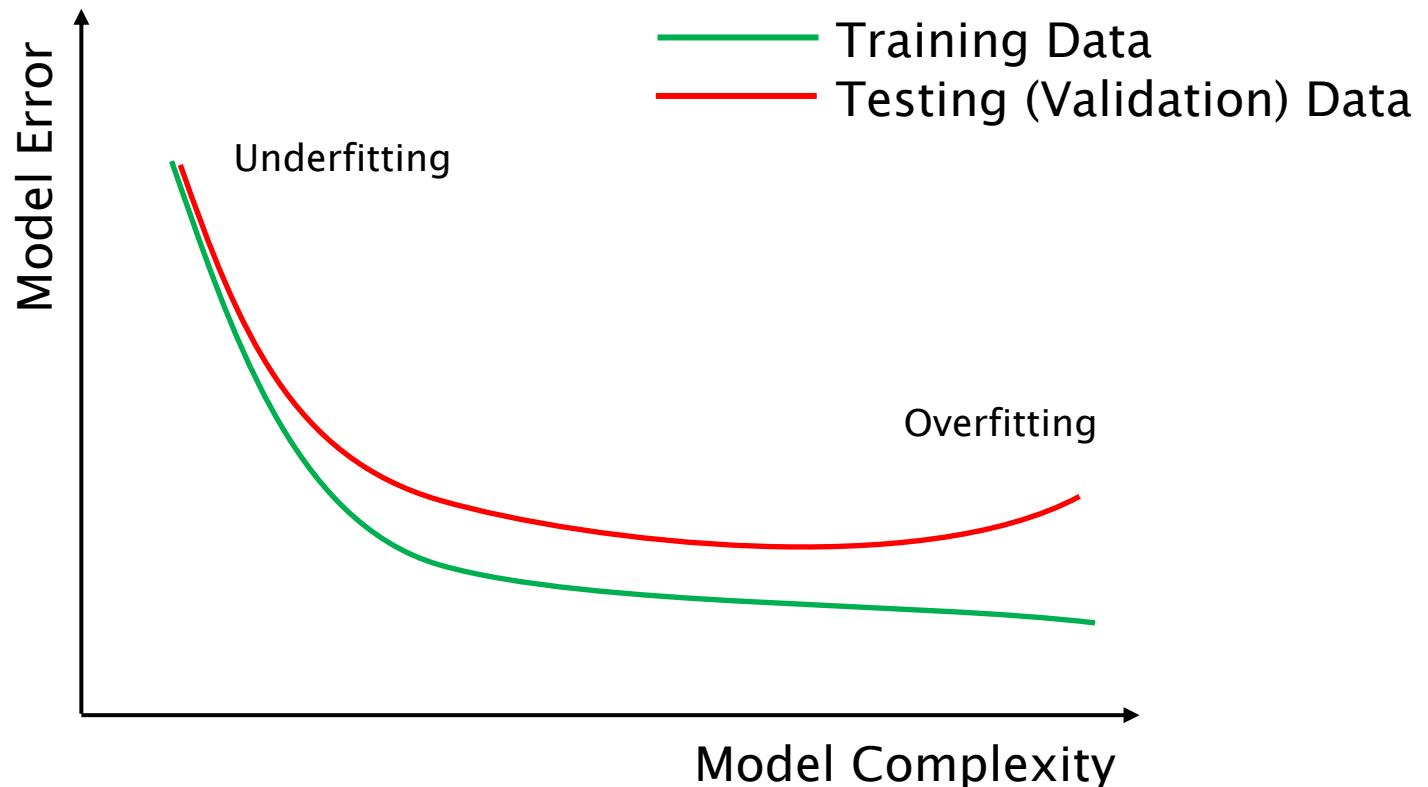
Underfitting vs. Overfitting



If A and B are noise, then h_2 overfits.

If A and B are *not* noise, then h_1 underfits.

Model (Hypothesis) Selection



Supervised Learning

Bayesian Learning
Neural Networks

Bayesian Learning

- ▶ Combines prior knowledge with evidence to make predictions
- ▶ Optimal classifier (not practical)
 - Need to know all true probabilities
- ▶ Naïve Bayes classifier (practical)
 - Assumes independence among features

Bayes Rule

$$P(C_i | \mathbf{x}) = \frac{P(\mathbf{x} | C_i)P(C_i)}{P(\mathbf{x})}$$



Thomas Bayes

- ▶ C_i is the class, $1 \leq i \leq K$
- ▶ \mathbf{x} is the feature values of an instance
- ▶ $P(C_i | \mathbf{x})$ = probability that instance \mathbf{x} belongs to class C_i (*posterior*)
- ▶ $P(\mathbf{x} | C_i)$ = probability that an instance drawn from class C_i would be \mathbf{x} (*likelihood*)
- ▶ $P(C_i)$ = probability of class C_i (*prior*)
- ▶ $P(\mathbf{x})$ = probability of instance \mathbf{x} (*evidence*)

Bayes Rule: Family Car

$$P(FamilyCar | EnginePower, Price) = \frac{P(EnginePower, Price | FamilyCar)P(FamilyCar)}{P(EnginePower, Price)}$$

Bayes Classifier

- ▶ Classify instance x as class C_i such that

$$i = \arg \max_{1 \leq k \leq K} P(C_k | x)$$

- ▶ Since only interested in maximum, can ignore denominator $p(x)$ (i.e., ignore normalization α)

$$i = \arg \max_{1 \leq k \leq K} P(x | C_k)P(C_k)$$

- ▶ If prior probability distribution of classes is uniform, then can ignore $P(C_i)$

$$i = \arg \max_{1 \leq k \leq K} P(x | C_k)$$

Bayes Classifier

- ▶ Practical issue #1
 - $P(x|C_i)$ is a joint probability distribution
 - Need to know the probability of every possible instance given every possible class
 - Even for D boolean attributes and K classes, that's K^*2^D probabilities
- ▶ Solution
 - Assume attributes are independent of each other

$$p(x_1, x_2, \dots, x_D | C_i) = \prod_{j=1}^D p(x_j | C_i)$$

Naïve Bayes Classifier

- Given training set X
- Estimate probabilities from X

$$P(C_i) = \frac{|\{(x,y) \in X \mid y = C_i\}|}{|X|}$$

$$P(x_j = v \mid C_i) = \frac{|\{(x,y) \in X \mid x_j = v \wedge y = C_i\}|}{|\{(x,y) \in X \mid y = C_i\}|}$$

- Classify new instance x as class C_i such that

$$i = \arg \max_{1 \leq k \leq K} P(C_k) * \prod_{j=1}^D P(x_j \mid C_k)$$

Example

- ▶ $P(\text{FamilyCar})$
 - $P(\text{FamilyCar=yes}) = 3/8$
 - $P(\text{FamilyCar=no}) =$
- ▶ $P(\text{Price} | \text{FamilyCar})$
 - $P(\text{Price}=10000 | \text{FamilyCar=yes}) = 2/3$
 - $P(\text{Price}=20000 | \text{FamilyCar=yes}) =$
 - $P(\text{Price}=30000 | \text{FamilyCar=yes}) =$
 - $P(\text{Price}=10000 | \text{FamilyCar=no}) =$
 - $P(\text{Price}=20000 | \text{FamilyCar=no}) =$
 - $P(\text{Price}=30000 | \text{FamilyCar=no}) =$
- ▶ $P(\text{EnginePower} | \text{FamilyCar})$
 - $P(\text{EnginePower}=100 | \text{FamilyCar=yes}) = 2/3$
 - $P(\text{EnginePower}=200 | \text{FamilyCar=yes}) =$
 - $P(\text{EnginePower}=300 | \text{FamilyCar=yes}) =$
 - $P(\text{EnginePower}=100 | \text{FamilyCar=no}) =$
 - $P(\text{EnginePower}=200 | \text{FamilyCar=no}) =$
 - $P(\text{EnginePower}=300 | \text{FamilyCar=no}) =$
- ▶ $P(\text{FamilyCar=yes} | \text{Price}=20000 \wedge \text{EnginePower}=200) = ?$

Price	EnginePower	FamilyCar
10000	100	yes
10000	200	yes
10000	300	no
20000	100	yes
20000	300	no
30000	100	no
30000	200	no
30000	300	no

Example (cont.)

- ▶ $P(\text{FamilyCar=yes} \mid \text{Price}=20000 \wedge \text{EnginePower}=200) = ?$

Naïve Bayes Classifier

► Practical issue #2

- What if no examples in class C_i have $x_j = v$?

$$P(x_j = v | C_i) = 0$$

$$P(C_i) * \prod_{j=1}^D P(x_j | C_i) = 0$$

► Solution

$$p(x_j = v | C_i) = \frac{|\{(x, y) \in X \mid x_j = v \text{ and } y = C_i\}| + 1}{|\{(x, y) \in X \mid y = C_i\}| + |\text{values}(x_j)|}$$

Another Example

- ▶ $P(\text{FamilyCar=yes} \mid \text{Price}=30000 \wedge \text{EnginePower}=200) = ?$
 $= \alpha * P(\text{Price}=30000 \mid \text{FamilyCar=yes}) *$
 $P(\text{EnginePower}=200 \mid \text{FamilyCar=yes}) *$
 $P(\text{FamilyCar=yes})$
 $= \alpha (0/3)(1/3)(3/8) = ?$
- ▶ $\text{values(Price)} = \{10000, 20000, 30000\}$
- ▶ $|\text{values(Price)}| = 3$
- ▶ $P(\text{Price}=30000 \mid \text{FamilyCar=yes}) = ?$

Naïve Bayes Classifier

- ▶ Practical issue #3
 - What if x_j is a continuous feature?
- ▶ Solution 1
 - Assume some parameterized distribution for x_j
 - E.g., normal
 - Learn parameters of distribution from data
 - E.g., mean and variance of x_j values
- ▶ Solution 2
 - Discretize feature
 - E.g., price $\in \mathbb{R}$ to price $\in \{\text{low}, \text{medium}, \text{high}\}$

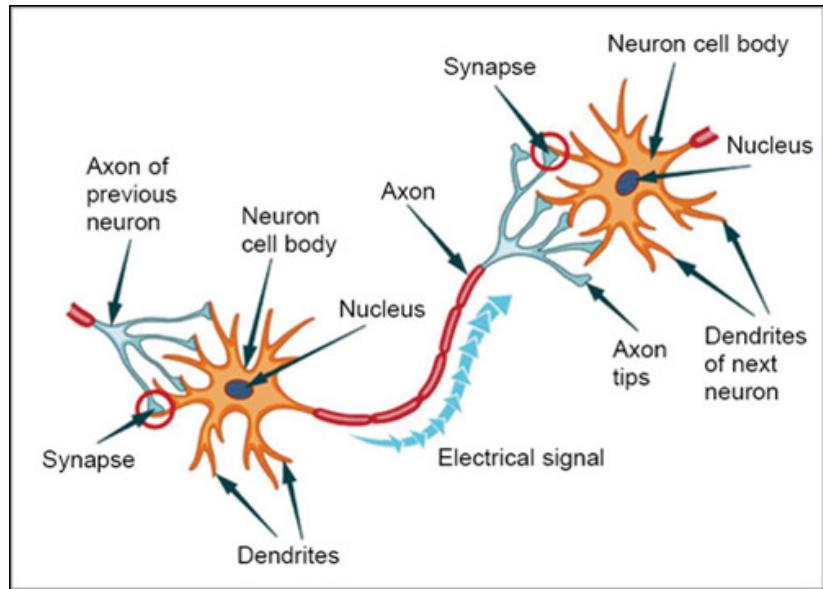
Bayesian Learning: Summary

- ▶ Optimal learner, in theory
- ▶ Naïve Bayes learner, practical
 - Independence assumption rarely true
 - E.g., Is “price” independent of “engine power”?
 - Naïve Bayes learner still does surprisingly well
 - Simple, effective baseline for other learners

Neural Networks

Neural Networks

- ▶ Inspired from brain
 - Brain consists of interconnected neurons
 - Humans still outperform machines in most areas of intelligence



Human Brain vs. Computer

▶ Processors

- Computer
 - CPUs: 20 cores (10^9 Hz)
 - GPUs: 10^4 cores (10^9 Hz)
- Human brain
 - 10^{11} neurons (10^3 Hz)

▶ Parallelism

- Computer
 - CPUs some, GPUs lots
 - Few interconnections
- Human brain
 - 10^{15} synapses
 - Each neuron has $\sim 10^4$ connections to other neurons



The Singularity

1 The accelerating pace of change ...



2 ... and exponential growth in computing power ...

Computer technology, shown here climbing dramatically by powers of 10, is now progressing more each hour than it did in its entire first 90 years.

COMPUTER RANKINGS

By calculations per second per \$1,000



Analytical engine
Never fully built, Charles Babbage's invention was designed to solve computational and logical problems



Colossus
The electronic computer, with 1,500 vacuum tubes, helped the British crack German codes during WW II

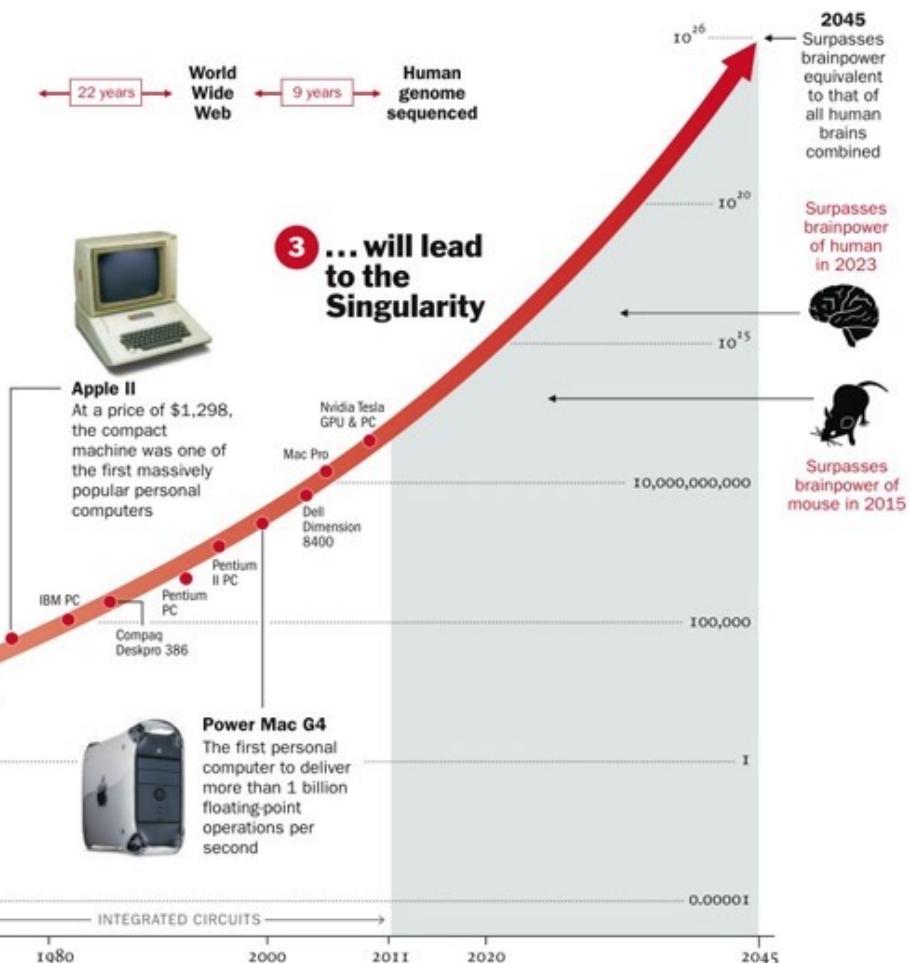


UNIVAC I
The first commercially marketed computer, used to tabulate the U.S. Census, occupied 943 cu. ft.



Apple II
At a price of \$1,298, the compact machine was one of the first massively popular personal computers

3 ... will lead to the Singularity



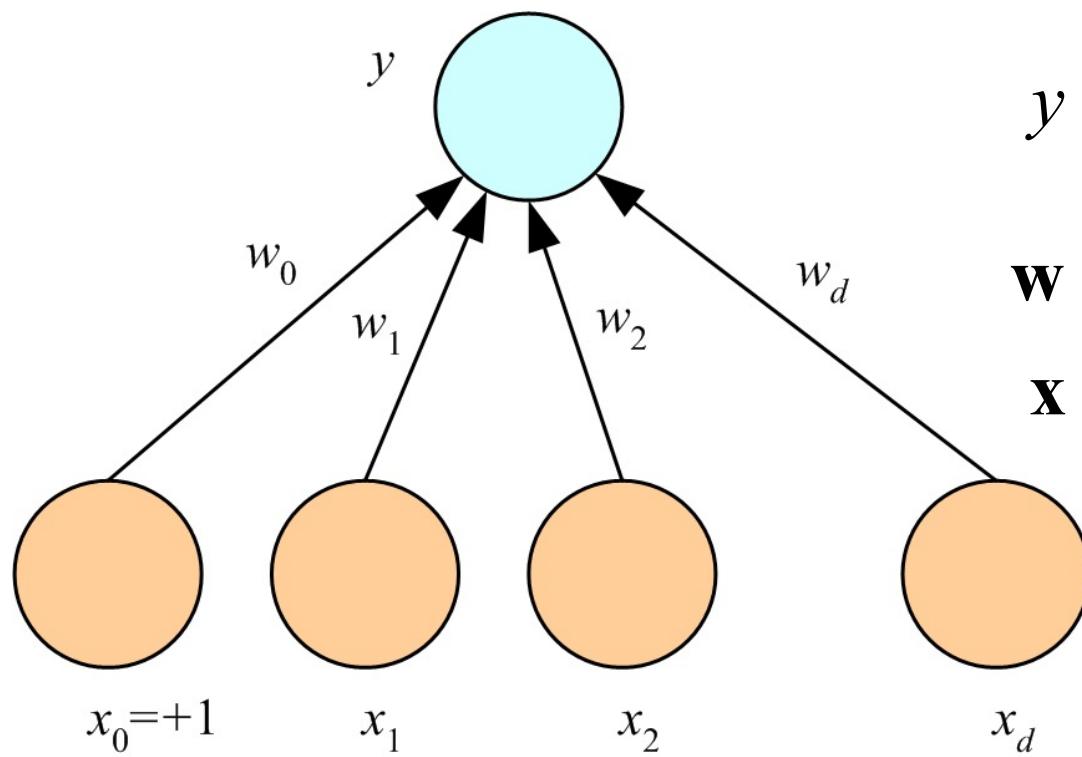
Time Magazine (Feb 2011)

What Happens Then...?



Terminator 2 Judgment Day (1991)

Perceptron (Rosenblatt, 1962)



$$y = \sum_{j=1}^d w_j x_j + w_0 = \mathbf{w} \cdot \mathbf{x}$$

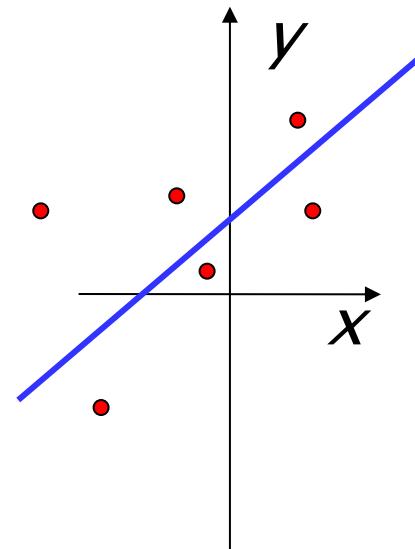
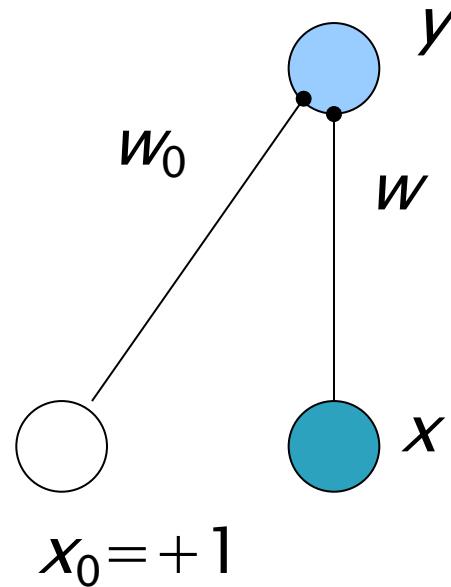
$$\mathbf{w} = [w_0, w_1, \dots, w_d]$$

$$\mathbf{x} = [1, x_1, \dots, x_d]$$

Learn function f
Modify weights so $y = f(x)$

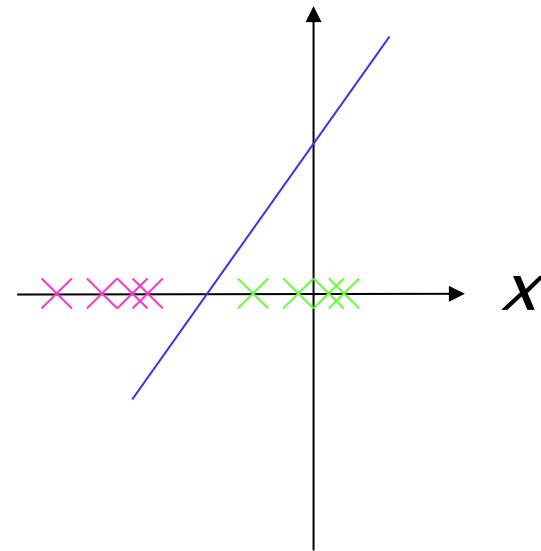
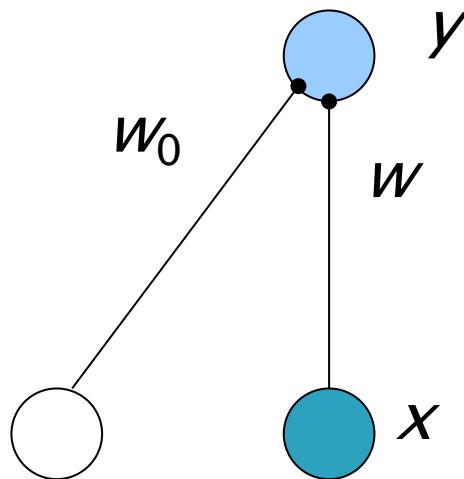
Perceptron Regression

► $y = wX + w_0$



Perceptron Classification

- If $(wx + w_0 > 0)$ Then $y=1$ Else $y=0$



Perceptron Training

- ▶ Change weights to reduce error
- ▶ Gradient descent

$$\Delta w_i = \eta(y_j - o_j)x_{ji}$$

- o_j = output of perceptron for example j
- Learning rate η controls rate of descent

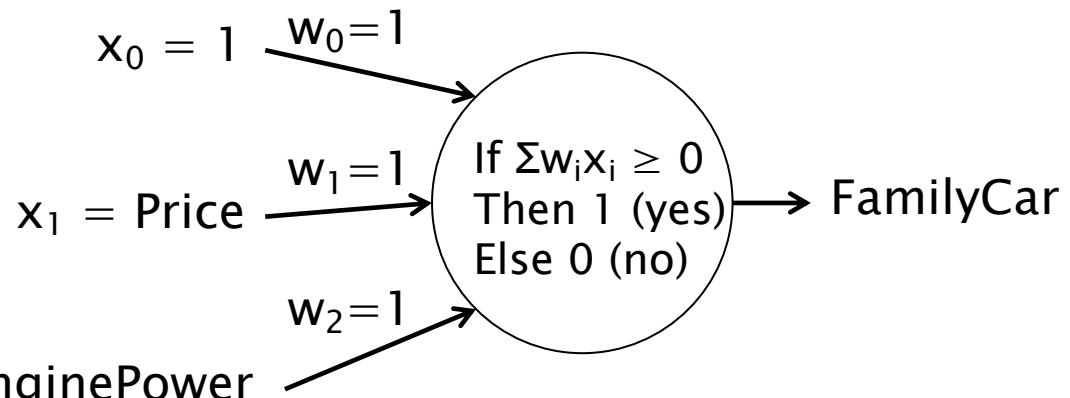
Example

Price	EnginePower	FamilyCar
1	1	1
1	2	1
1	3	0
2	1	1
2	3	0
3	1	0
3	2	0
3	3	0

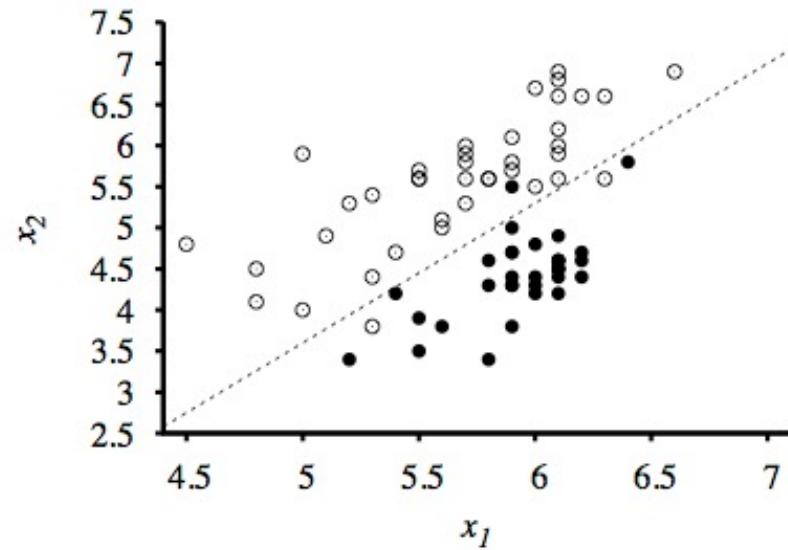
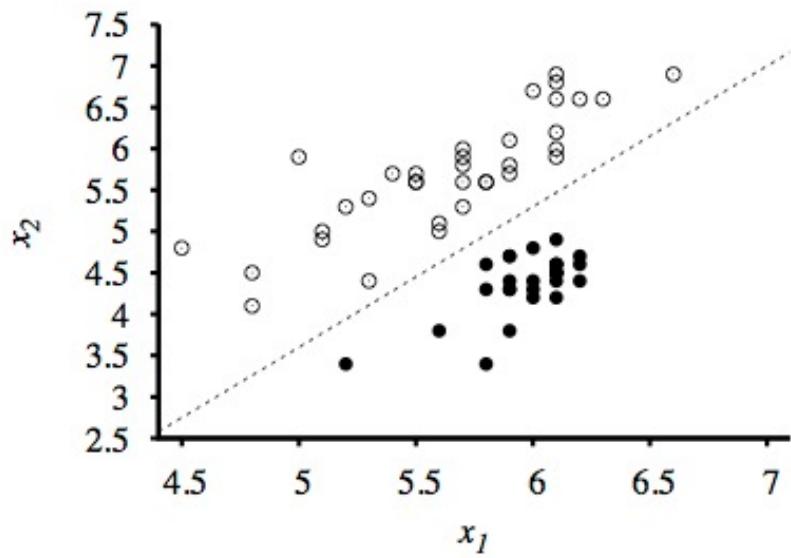
$x_2 = \text{EnginePower}$

$$\Delta w_i = \eta(y_j - o_j)x_{ji}$$

Let $\eta = 0.5$

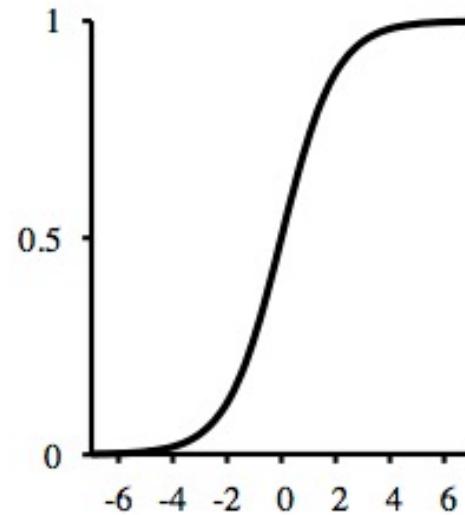
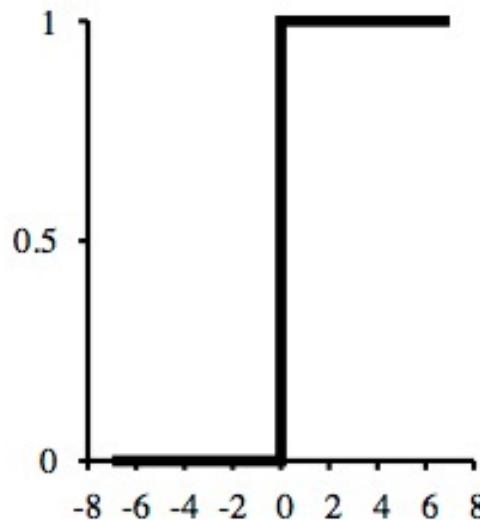


Linearly Separable



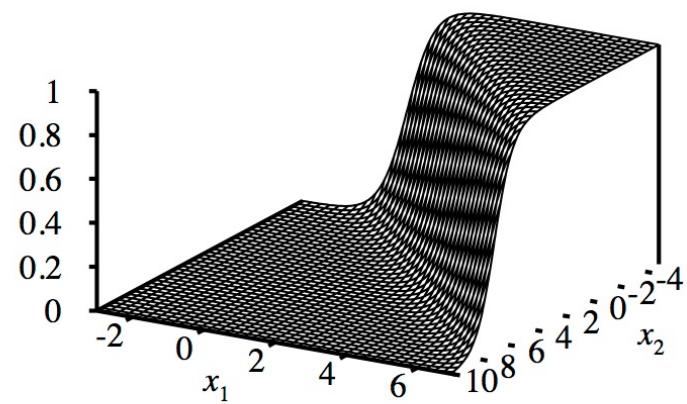
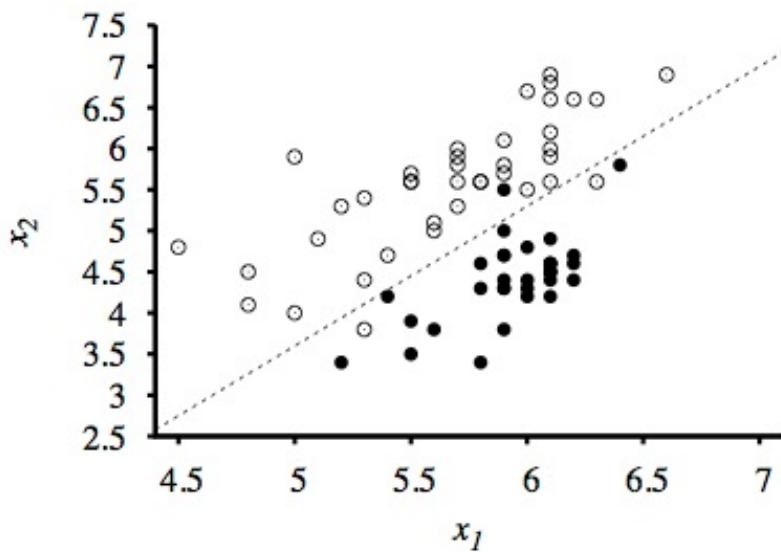
Logistic Function

- ▶ Also called “sigmoid function”

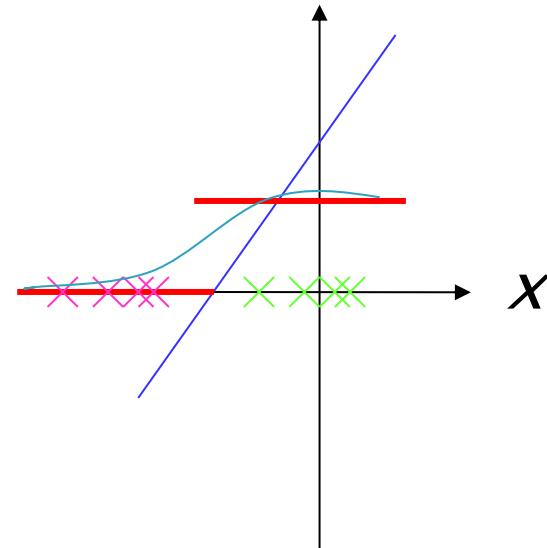
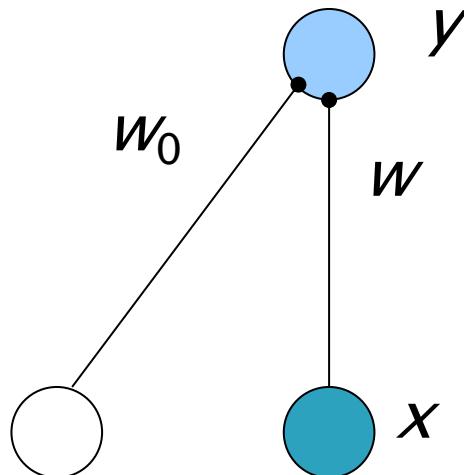


$$\text{logistic}(z) = \frac{1}{1 + e^{-z}}$$

Logistic Function



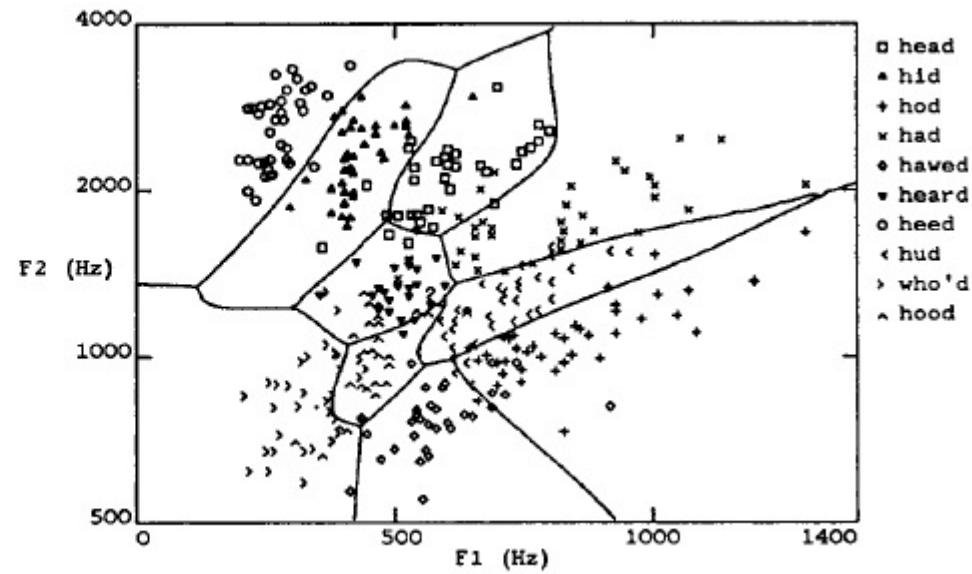
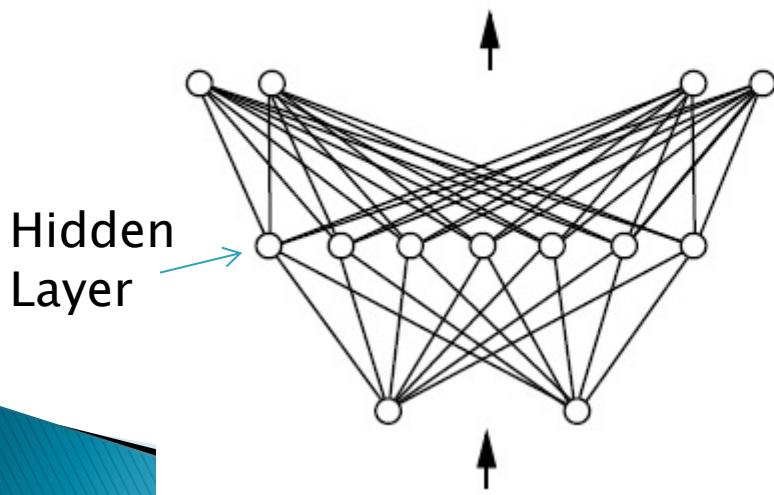
Perceptron Classification



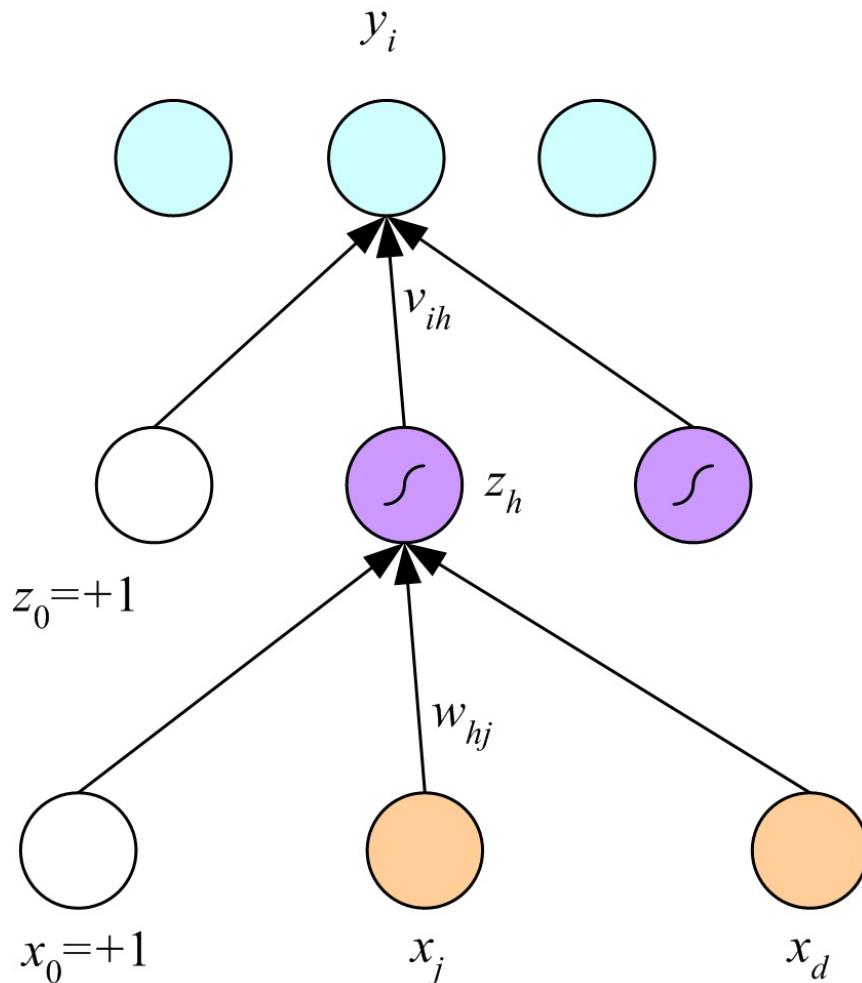
$$y = \text{logistic}(\mathbf{w} \cdot \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$

Learning Nonlinear Functions

- ▶ Perceptrons can only approximate linear functions
- ▶ But multiple layers of logistic perceptrons can approximate any nonlinear function



Multilayer Perceptrons



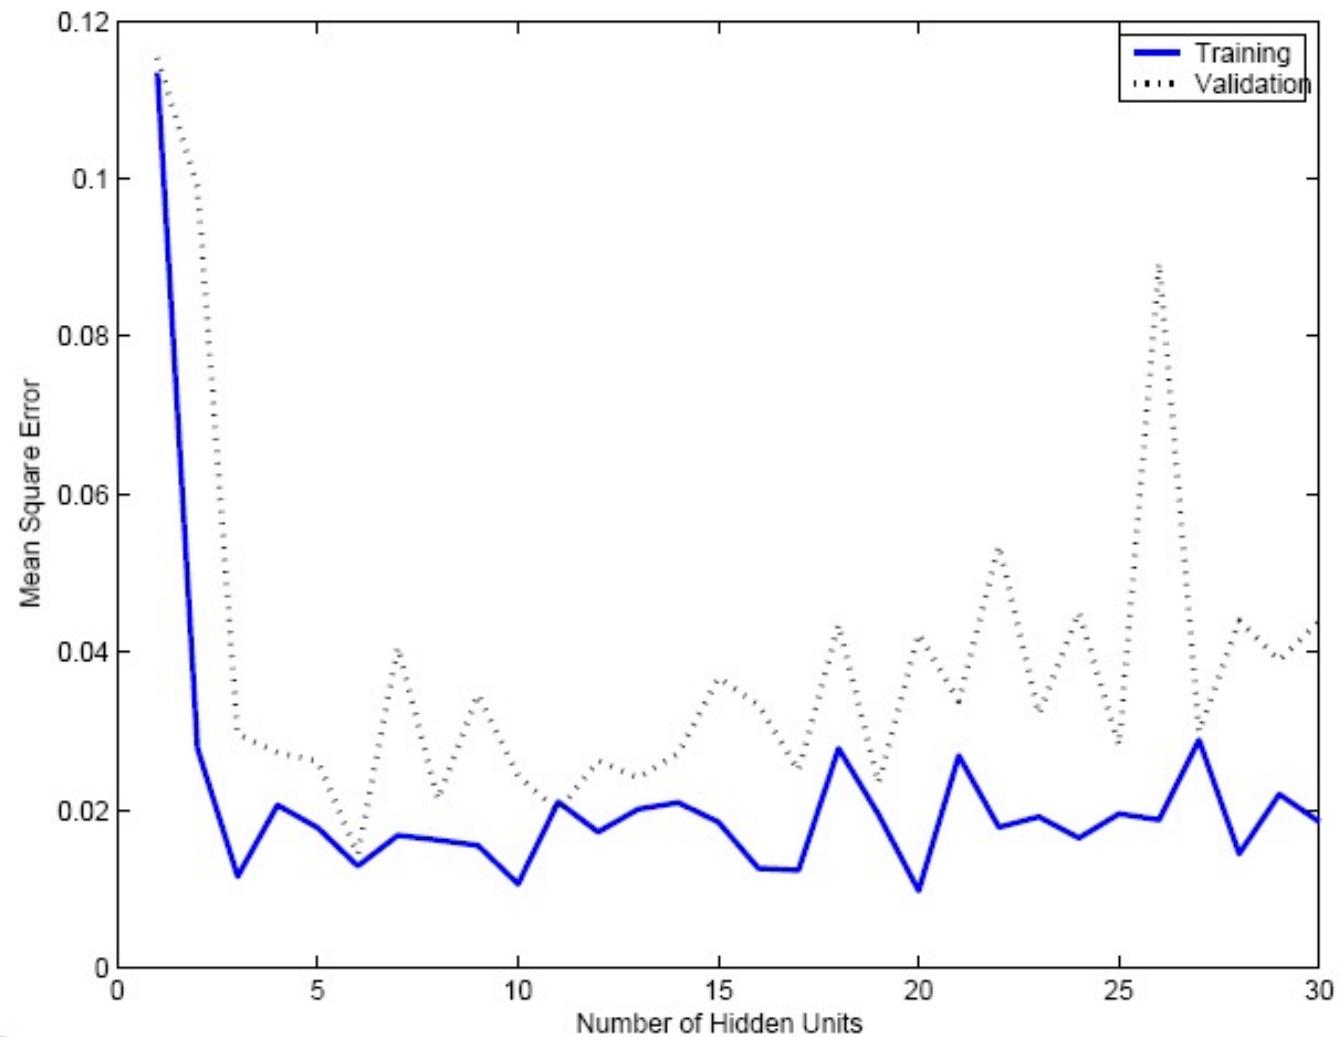
$$y_i = \mathbf{v}_i \cdot \mathbf{z} = \sum_{h=1}^H v_{ih} z_h + v_{i0}$$

$$\begin{aligned} z_h &= \text{logistic}(\mathbf{w}_h \cdot \mathbf{x}) \\ &= \frac{1}{1 + \exp\left[-\left(\sum_{j=1}^d w_{hj} x_j + w_{h0}\right)\right]} \end{aligned}$$

Learn weights by error backpropagation.
(Rumelhart et al., 1986)

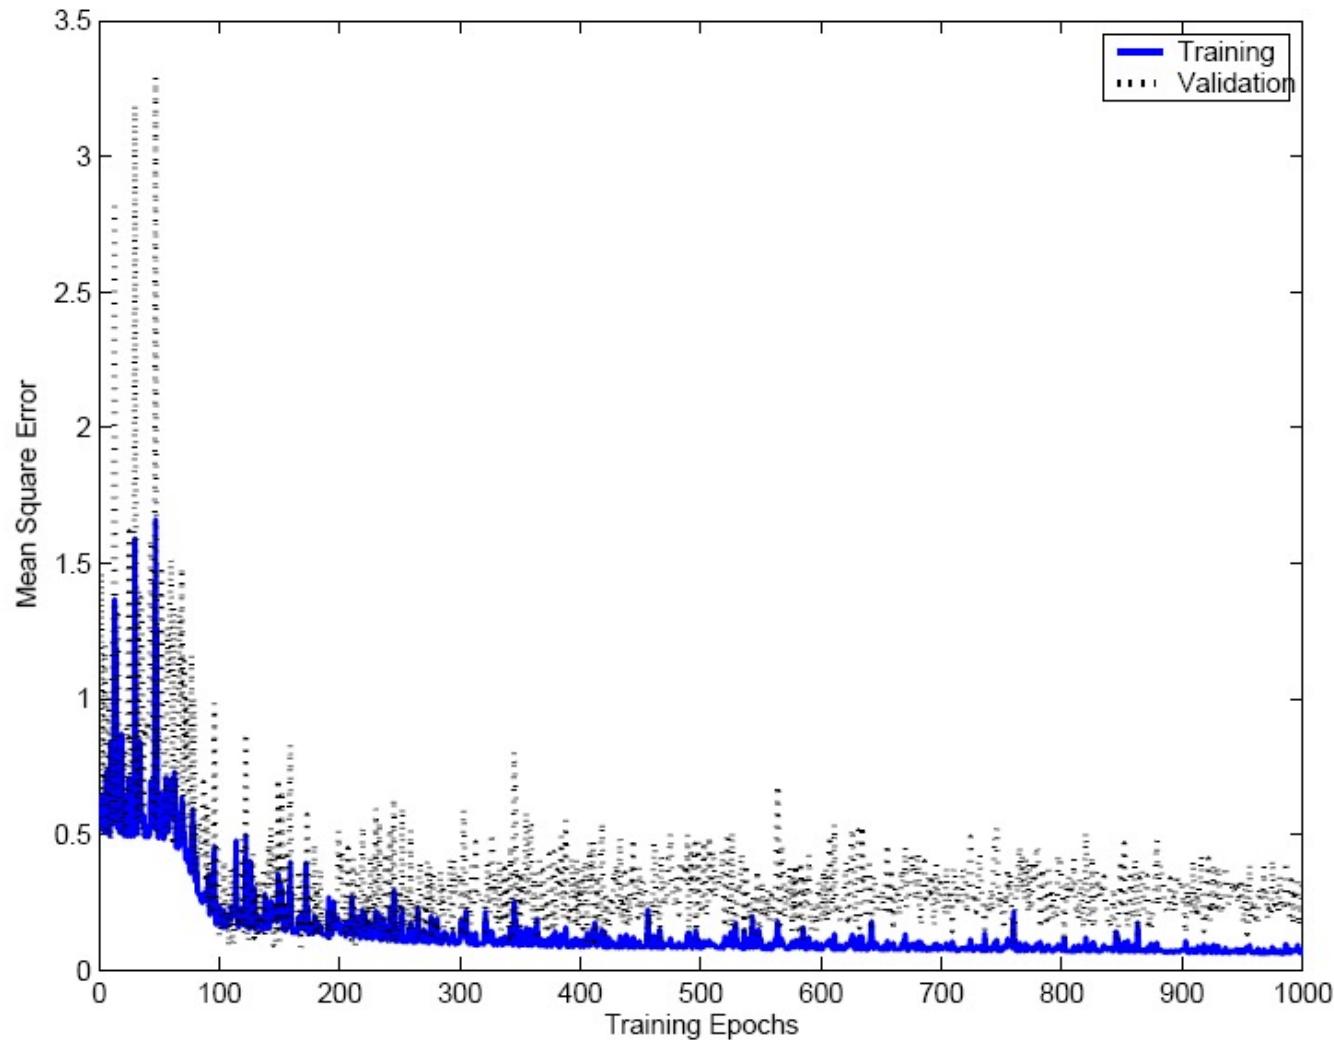
Overfitting in MLPs

$$f(x) = \sin(6x)$$



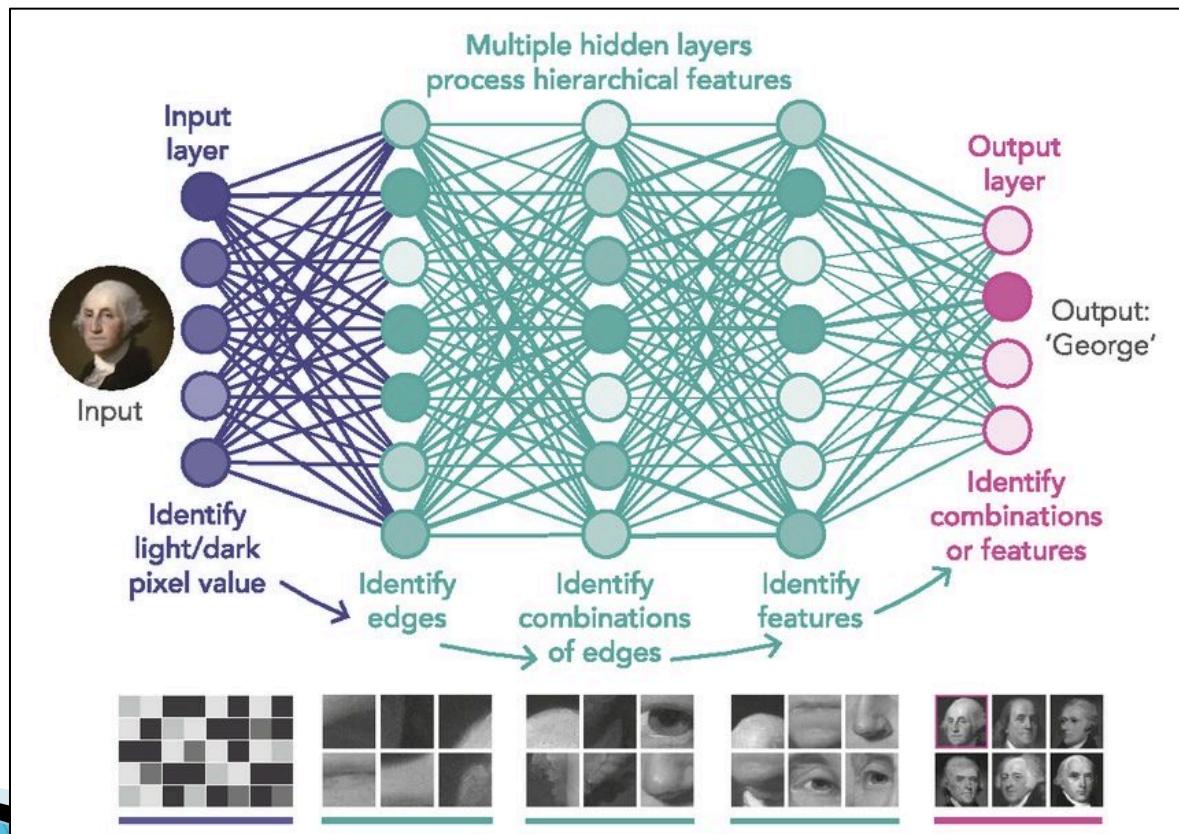
Overfitting in MLPs

- ▶ Similar overfitting behavior if training continued too long
- ▶ More and more weights move from zero
- ▶ Overtraining

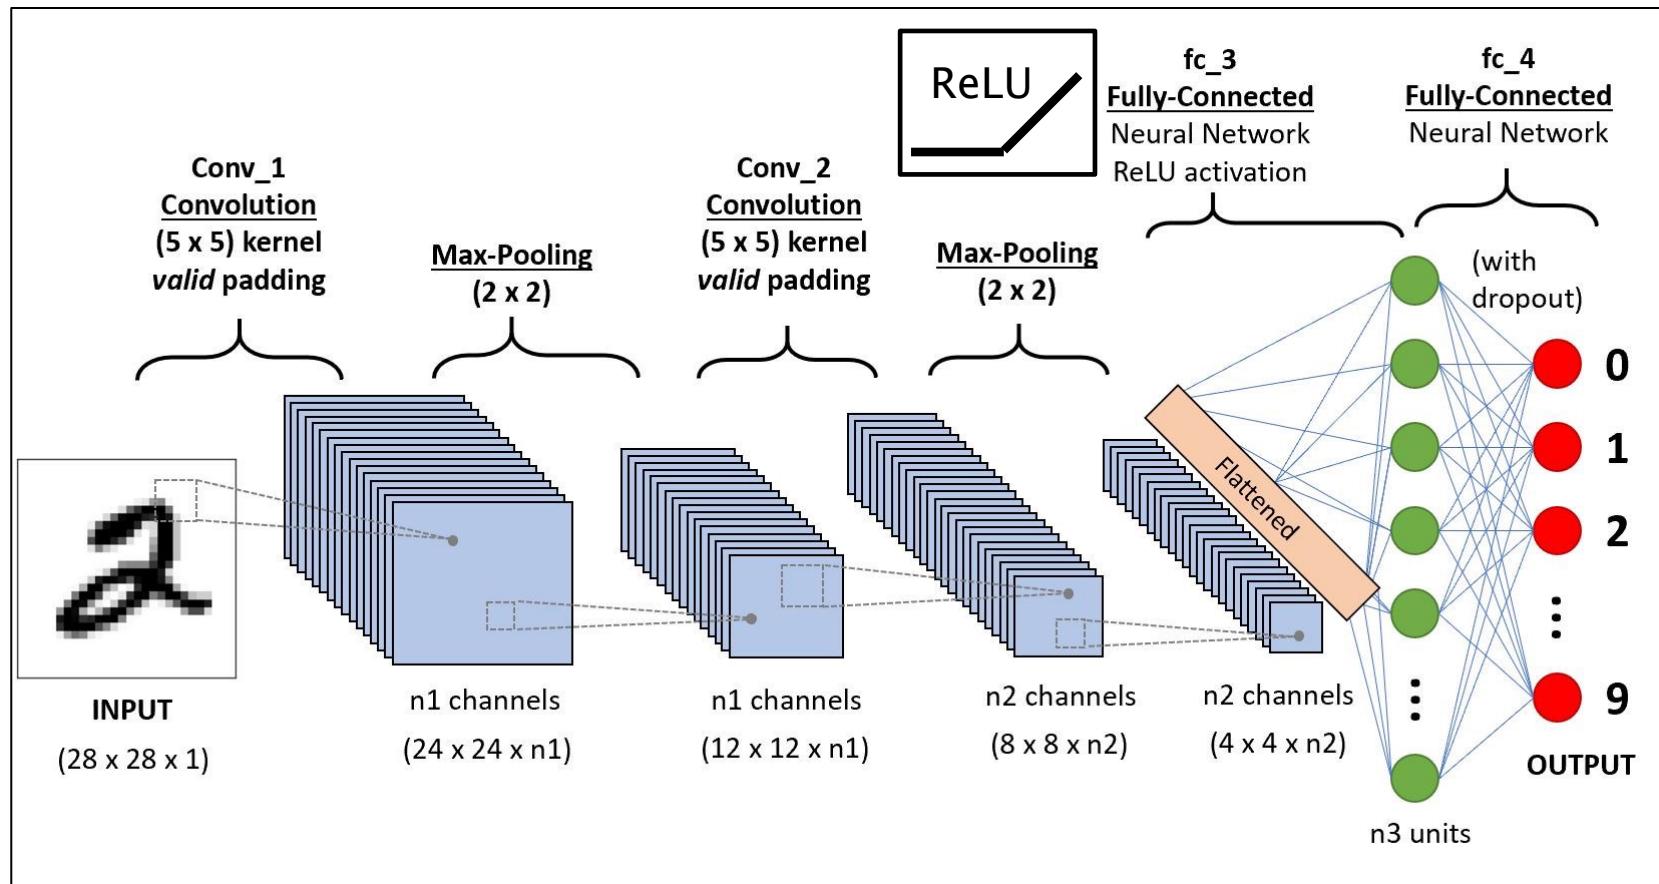


Deep Learning

- ▶ Using deep (many-layered) neural nets to learn complex abstractions (features)



Deep Convolutional Neural Net (CNN) for Image Classification

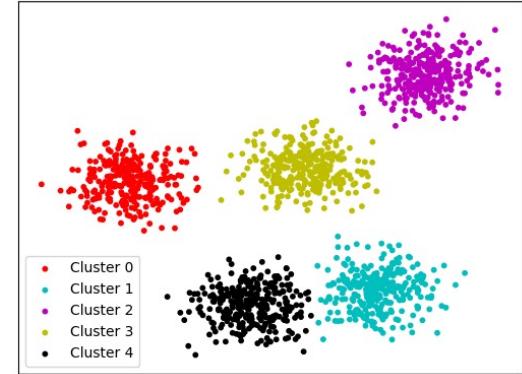


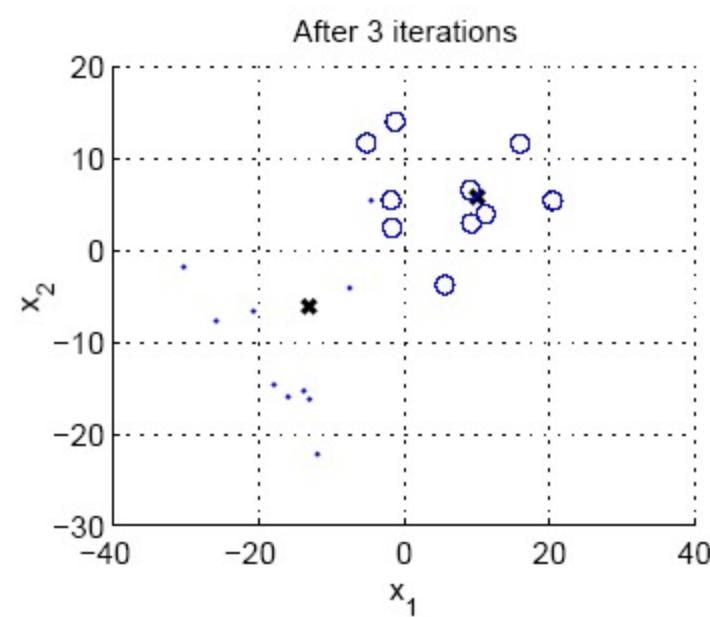
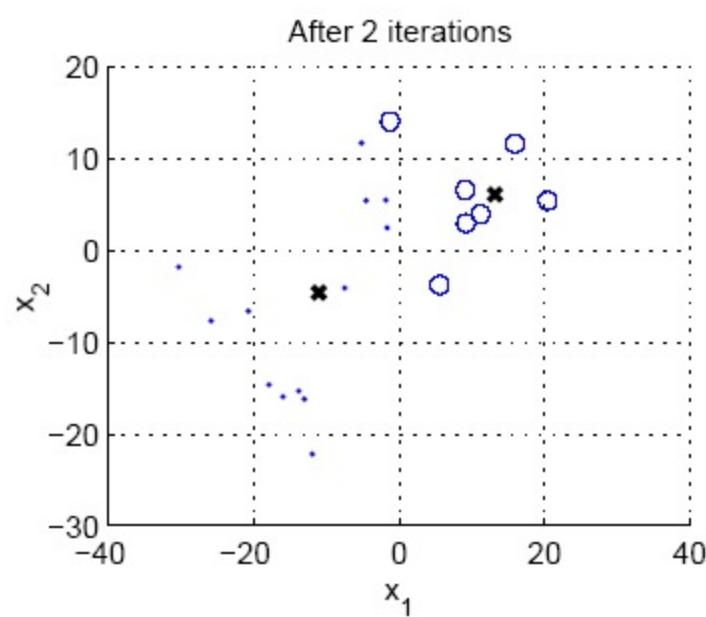
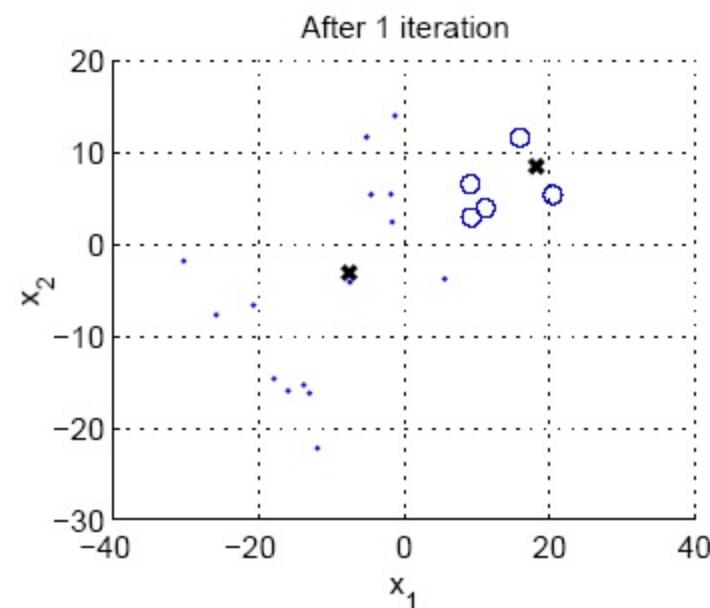
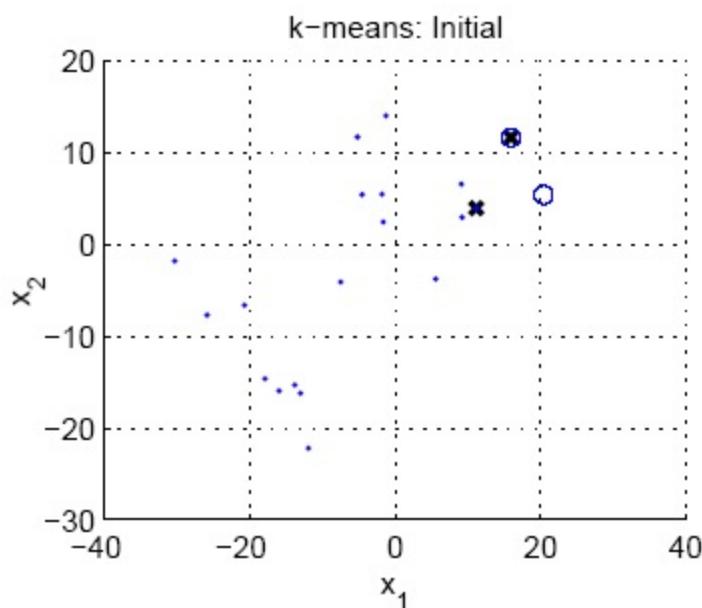
Unsupervised Learning

k-Means Clustering

k -means Clustering

- ▶ Unsupervised learning
 - Just feature values, no classes
- ▶ Partition instances into k disjoint sets
- ▶ Each set has a representative instance m_i
- ▶ Place instance x into set i such that $distance(x, m_i)$ is minimal
- ▶ Choose new central m_i for each set
- ▶ Repeat until m_i converge

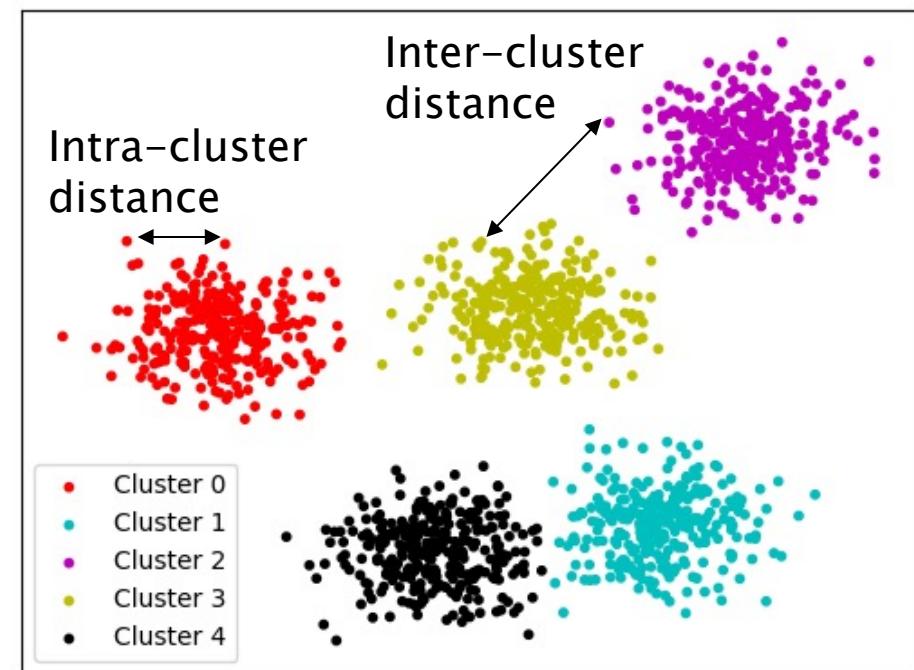




k-means Clustering

▶ Choosing k ?

- Try several ($2 \leq k \leq N$)
- Choose k minimizing intra-cluster distance and maximizing inter-cluster distance



ML Software Tools

- ▶ Waikato Environment for Knowledge Analysis (WEKA)
 - Java based
 - <http://www.cs.waikato.ac.nz/ml/weka>
- ▶ SciKit Learn
 - Python based
 - <http://scikit-learn.org>
- ▶ Deep Learning (mostly Python)
 - TensorFlow (<https://www.tensorflow.org>)
 - Keras (<https://keras.io>)

Reinforcement Learning

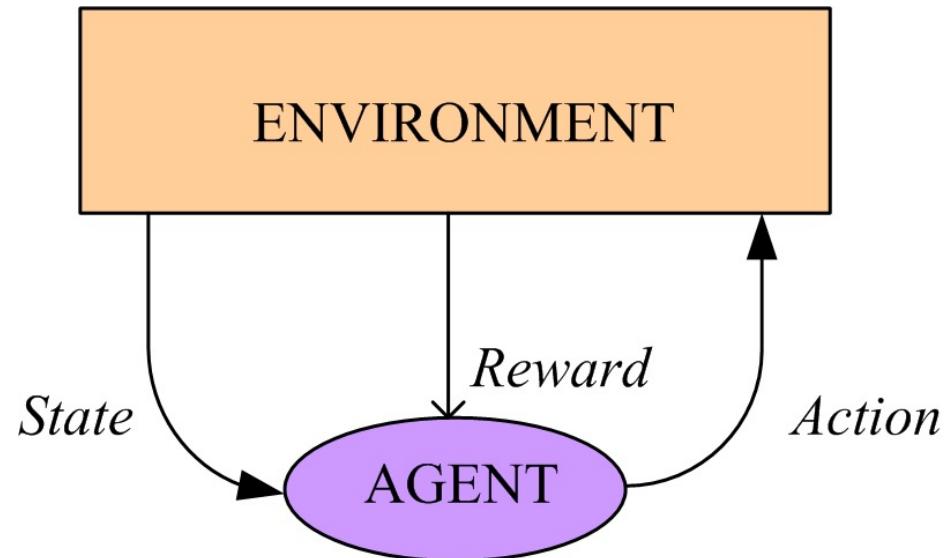
Reinforcement Learning



sarvagyavaish.github.io/FlappyBirdRL

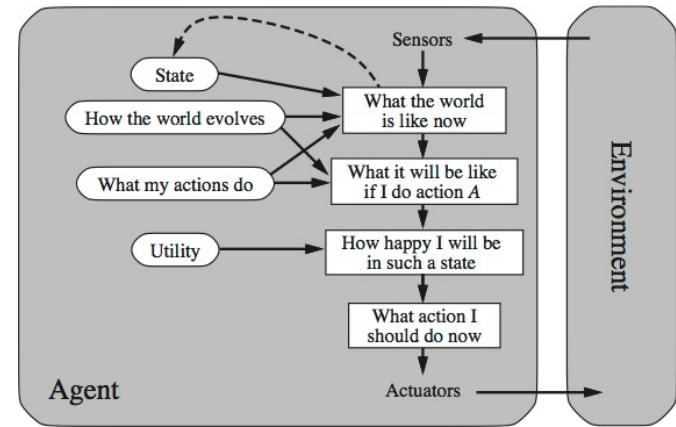
Reinforcement Learning

- ▶ Agent in some state in the environment, takes an action and sometimes receives reward, and the state changes
- ▶ Delayed reward
- ▶ Credit-assignment
- ▶ Learn a policy
 - π : State \rightarrow Action
- ▶ Applications
 - Game-playing
 - Robot control



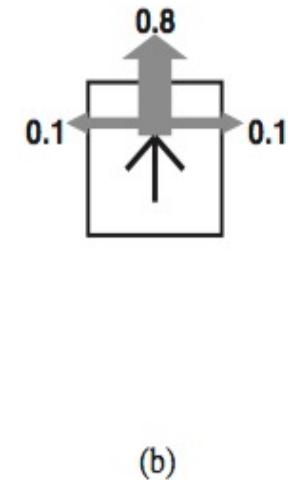
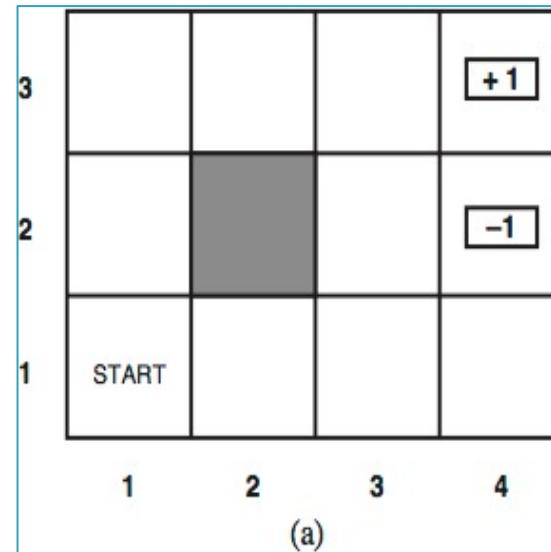
Reinforcement Learning Agents

- ▶ Utility-based agent
 - Learn utility function on states
 - Choose actions to maximize expected utility
 - Requires action model
 - Action: state → state
- ▶ Q-learning agent
 - Learns Q function: value of taking action A in state S
 - No action model, but learning slower
- ▶ Passive learning of utility or Q functions
 - Policy given
- ▶ Active learning: Also learn policy
 - Exploration



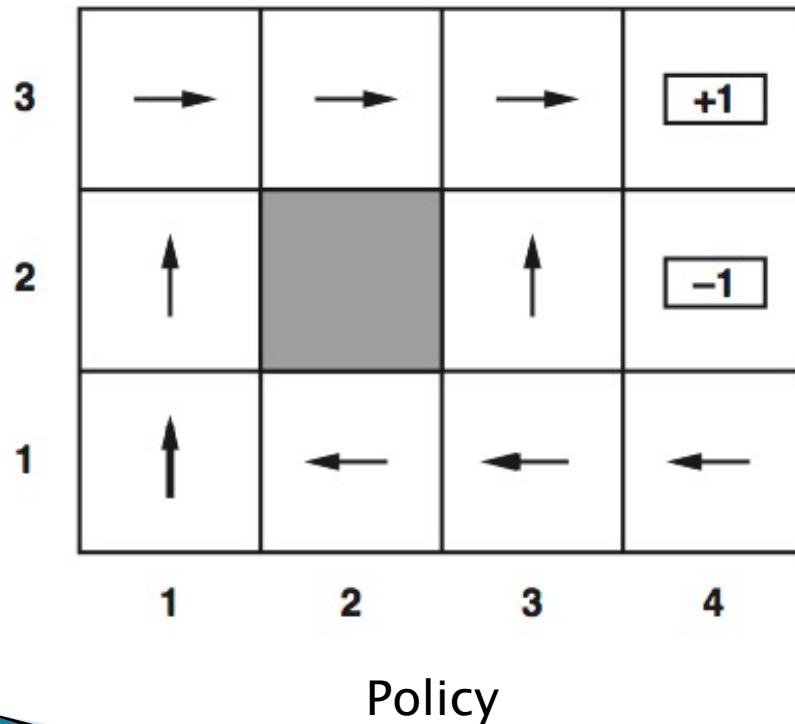
Simple Environment

- ▶ Actions: Up, Down, Left, Right
- ▶ 80% chance intended move executed
- ▶ 20% chance agent moves 90 degrees from intended direction
- ▶ Agent bumps into walls and obstacle
- ▶ Reward
 - +1 for reaching (4,3)
 - -1 for reaching (4,2)
 - -0.04 for other moves
- ▶ Fully observable



Passive Reinforcement Learning

- Given policy $\pi(s)$
- Learn utility function $U^\pi(s)$



	1	2	3	4
3	0.812	0.868	0.918	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388

Utilities

Utility Estimation

- ▶ Utility of state s is the reward in state s plus the expected utility of its successor states

$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) U^\pi(s')$$

- ▶ γ is the discount factor ($0 < \gamma \leq 1$)

Utility Estimation: Example

- ▶ Let $\gamma=1$
- ▶ Using previous example and policy
 - Initially all $U(s)=0$
 - Except $U([4,3])=1$ and $U([4,2])=-1$
- ▶ $U([3,3]) = ?$, policy says go Right
- ▶ Eventually, $U([3,3])$
 - $= (-0.04) + (0.8)*(1) + (0.1)*(0.918) + (0.1)*(0.660)$
 - $= 0.918$

Temporal Difference Learning

- ▶ Observe transition from s to s'
- ▶ Adjust utility of state s to be closer to utility of state s'

$$U^\pi(s) = U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

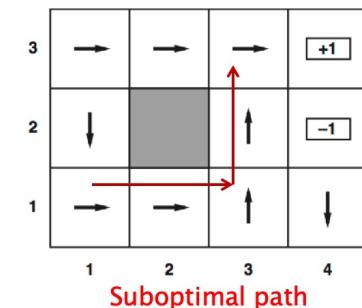
- ▶ Use learning rate α to control rate of adjustment
- ▶ If α decreases with the number of times state s is visited, then TD will converge to correct $U^\pi(s)$

Active Reinforcement Learning

- ▶ Learn policy $\pi(s)$
- ▶ Agent in state s takes action maximizing expected utility

$$\pi(s) = \arg \max_a \sum_{s'} P(s'|s, a) * U(s')$$

- ▶ Initial actions random
- ▶ Estimate $P(s' | s, a)$ and $U(s)$ over time
- ▶ Can result in suboptimal policy



Exploration

- ▶ Force agent to explore unvisited states to avoid suboptimal policies
- ▶ Exploration function

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

- u is the utility estimate so far
- n is the number of times state s has been visited
- R^+ is estimate of best possible reward
- N_e is the number of times an action-state must be tried before relying on utility estimate

Q-Learning

- ▶ Active temporal-difference learning agent
- ▶ Define $Q(s,a)$ as the value of taking action \underline{a} in state \underline{s}

$$U(s) = \max_a Q(s, a)$$

$$Q(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q(s', a')$$

- ▶ Q function combines $\pi(s)$, $P(s'|s,a)$ and $U(s)$

Q-Learning

- ▶ Use TD approach to eliminate $P(s'|s,a)$

$$Q(s,a) \leftarrow Q(s,a) + \alpha(R(s) + \gamma \max_{a'} Q(s',a') - Q(s,a))$$

- ▶ Calculated when observe transition from s to s' using action a

Q-Learning Agent

function Q-LEARNING-AGENT (*percept*) **returns** an action

inputs: *percept*, a percept indicating current state s' and reward signal r'

persistent: Q , table of action values indexed by state and action, initially zero

N_{sa} , table of frequencies for state-action pairs, initially zero

s, a, r , previous state, action and reward, initially null

if TERMINAL?(s) **then**

foreach $a' \in \text{ACTIONS}(s)$

$Q[s, a'] \leftarrow r'$

if s is not null **then**

 increment $N_{sa}[s, a]$

$Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

$s, a, r \leftarrow s', \text{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a']), r'$

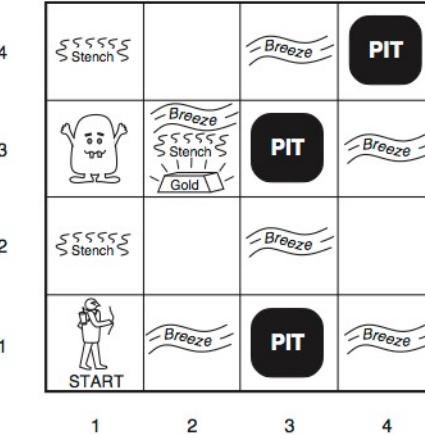
return a

Exploration
function $f(u, n)$

Learning rate (decreasing with n):
 $\alpha(n) = C / (C + n - 1)$

Wumpus World

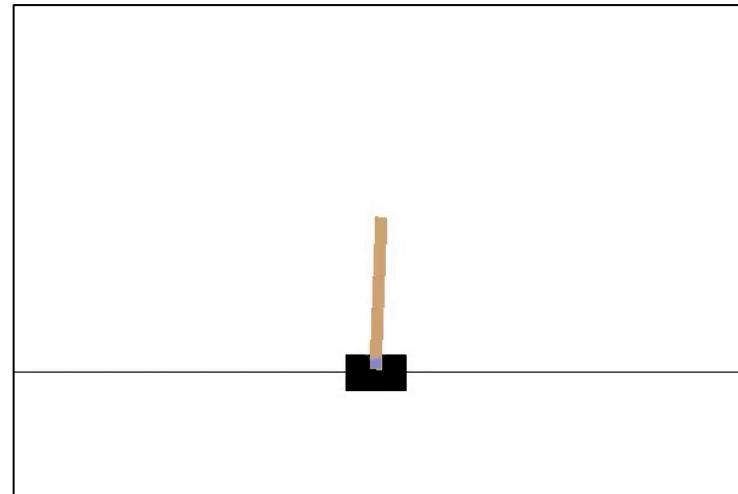
- ▶ Q-learning agent
- ▶ Limited to 4x4 worlds
- ▶ 512 states
 - 16 locations, 4 orientations, hasGold, hasArrow, inCave
- ▶ 6 actions
- ▶ Parameters
 - Best possible reward $R^+ = 1000$
 - Minimum state-action occurrences $N_e = 5$
 - Discount factor $\gamma = 0.9$
 - Learning rate $\alpha(n) = 100 / (99 + n)$



RL Software Tools

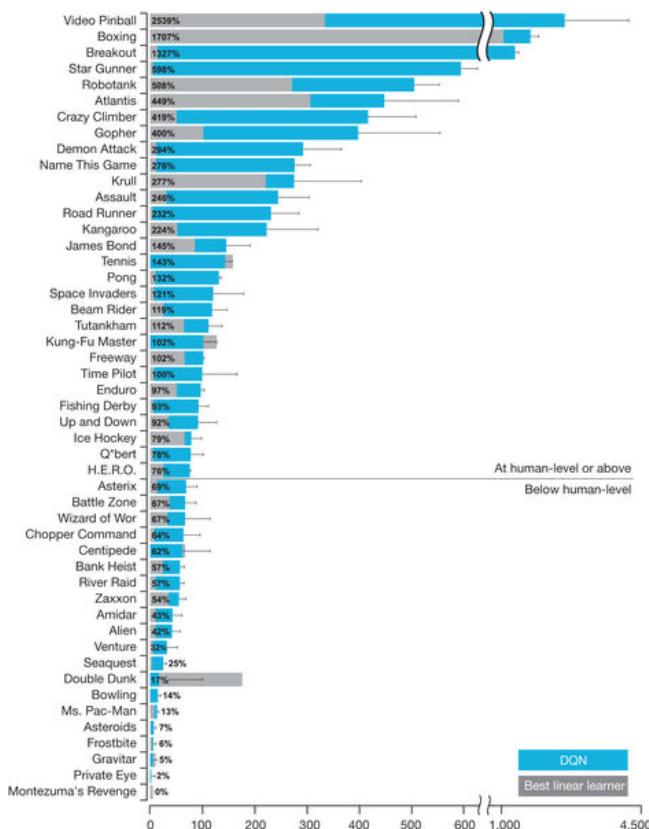
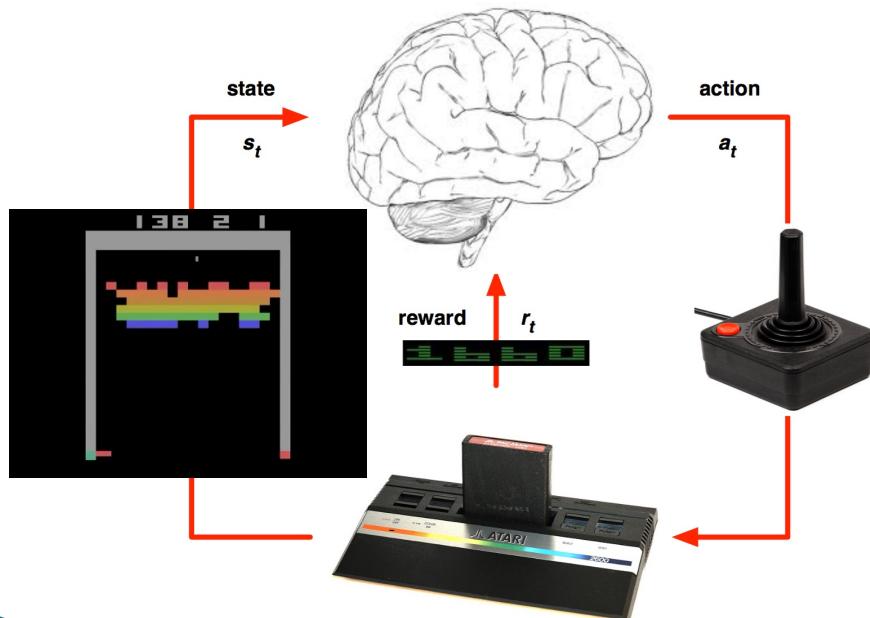
- ▶ OpenAI Gym
 - RL environments
 - gym.openai.com

Cart Pole



Deep Reinforcement Learning

- ▶ Deep Q-Learning Network (DQN)
 - Use Deep Learning Network to learn Q function
 - Agent “replays” past games



Summary: Learning

- ▶ Improving performance at some task through experience
- ▶ Supervised learning methods
 - Naïve Bayes
 - Neural network
- ▶ How to choose the right model?
 - Overfitting
- ▶ Unsupervised learning
 - Clustering
- ▶ Reinforcement learning

